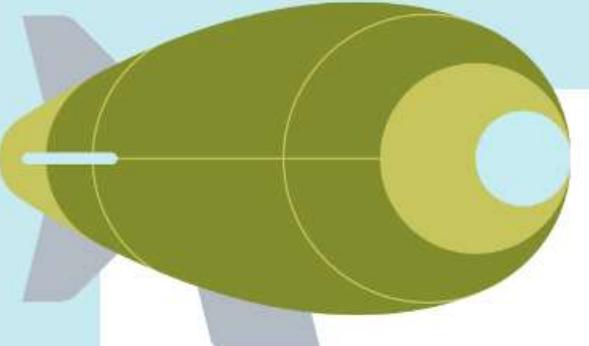


Database of the future is here - Azure SQL Hyperscale deep dive

Aditya Badramraju, Arvind Shyamsundar
Principal Product Managers, Microsoft Azure SQL DB



About us



Arvind Shyamsundar

Principal Product Manager, Microsoft

@arvisam (Bluesky, Mastodon and X)



Aditya Badramraju

Principal Product Manager, Microsoft

@aditya_feb22 (X)



Agenda

- Hyperscale **architecture**: motivations, approaches and **implementation**.
- **Working** with Hypescale databases.
- **Performance** monitoring and tuning.
- **Demos** to bring things to life!

Success criteria for us: leave you convinced that Hypescale is the best cloud native database for you!

<https://sqlb.it/?12734>



Related sessions

 More for less: Cost optimizing your Azure SQL databases – Friday, 1:40PM

 Experts Lounge: “Meet The PG: All things Azure SQL DB” – Friday, 12PM

 Microsoft booth in the Expo floor: Wednesday, 6-8PM; Thursday, 10:30AM-1PM
Saturday, 9-11AM

 “Perfecting business continuity for Azure SQL DB” – Saturday, 10:10AM

 “SQL DB: a developer's catalyst” – happening now ☺

 “A Deep Dive into DevOps Practices with Azure SQL” – Thursday, 10:10AM

Are you...



A developer?

Currently using Azure SQL DB?

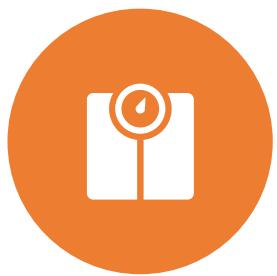
Long time SQL Server user, wanting to know more about Hyperscale?

Using other cloud DBs?

Using other DBs but mostly on-prem?



Challenges with databases



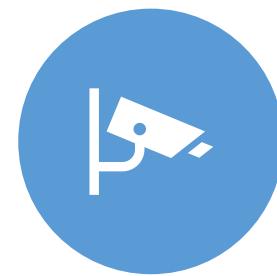
SCALE



PERFORMANCE



AVAILABILITY



SECURITY

What does “Hyperscale” mean?

Hyperscale computing

文 A 2 languages ▾

Article Talk

Read Edit View history Tools ▾

From Wikipedia, the free encyclopedia

(Redirected from [Hyperscale](#))

In computing, **hyperscale** is the ability of an [architecture](#) to [scale appropriately](#) as increased demand is added to the system.

This typically involves the ability to [seamlessly provide and add compute, memory, networking, and storage resources to a given node or set of nodes](#) that make up a larger [computing](#), [distributed computing](#), or [grid computing](#) environment. Hyperscale computing is necessary in order to build a robust and scalable [cloud](#), [big data](#), [map reduce](#), or [distributed storage](#) system and is often associated with the infrastructure required to [run large distributed](#) sites such as [Google](#), [Facebook](#), [Twitter](#), [Amazon](#), [Microsoft](#), [IBM Cloud](#) or [Oracle](#). Companies like [Ericsson](#), [AMD](#), and [Intel](#) provide hyperscale

Hyperscale in Azure SQL



Hyperscale (<https://aka.ms/hs>) is the trusted and well-known SQL Server engine, rearchitected for the cloud

- ✔ Distributed architecture, with independent services for storage and compute
- ✔ Scalable from 10 GB to 100TB and from 1 core to 128 vCore
- ✔ Multi-tiered cache architecture for great performance, no matter the database size
- ✔ Scalable (out) up to 30 secondary read-only replicas

“The only cloud relational database you will need”: use Hyperscale for all new applications on Azure SQL DB, no matter the size.

Architecting for scale-out



Shared Nothing

Log Shipping (Stand By Replica)
Always On availability groups



Shared Storage

Failover clusters (but then, that's not really scale out...)



Architecture “S”

Virtually unlimited read scale without data duplication
Storage capacity not being bounded by any single machine

Socrates: The New SQL Server in the Cloud

Panagiotis Antonopoulos, Alex Budovski, Cristian Diaconu, Alejandro Hernandez Saenz, Jack Hu, Hanuma Kodavalli, Donald Kossmann, Sandeep Lingam, Umar Farooq Minhas, Naveen Prakash, Vijendra Purohit, Hugh Qu, Chaitanya Sreenivas Ravella, Krystyna Reisterter, Sheetal Shrotri, Dixin Tang, Vikram Wakade
Microsoft Azure & Microsoft Research

ABSTRACT

The database-as-a-service paradigm in the cloud (DBaaS) is becoming increasingly popular. Organizations adopt this paradigm because they expect higher security, higher availability, and lower and more flexible cost with high performance. It has become clear, however, that these expectations cannot be met in the cloud with the traditional, monolithic database architecture. This paper presents a novel DBaaS architecture, called Socrates. Socrates has been implemented in Microsoft SQL Server and is available in Azure as SQL DB Hyperscale. This paper describes the key ideas and features of Socrates, and it compares the performance of Socrates with the previous SQL DB offering in Azure.

CCS CONCEPTS

• Information systems → DBMS engine architectures;

KEYWORDS

Database as a Service, Cloud Database Architecture, High Availability

ACM Reference Format:

Panagiotis Antonopoulos, Alex Budovski, Cristian Diaconu, Alejandro Hernandez Saenz, Jack Hu, Hanuma Kodavalli, Donald Kossmann, Sandeep Lingam, Umar Farooq Minhas, Naveen Prakash, Vijendra Purohit, Hugh Qu, Chaitanya Sreenivas Ravella, Krystyna Reisterter, Sheetal Shrotri, Dixin Tang, Vikram Wakade. 2019. Socrates: The New SQL Server in the Cloud. In *2019 International Conference on Management of Data (SIGMOD '19)*, June 30–July 5, 2019, Amsterdam, Netherlands. ACM, New York, NY, USA, 14 pages. <https://doi.org/10.1145/3299869.3314047>

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

SIGMOD '19, June 30–July 5, 2019, Amsterdam, Netherlands
© 2019 Association for Computing Machinery.
ACM ISBN 978-1-4503-5643-5/19/06...\$15.00
<https://doi.org/10.1145/3299869.3314047>

1 INTRODUCTION

The cloud is here to stay. Most start-ups are cloud-native. Furthermore, many large enterprises are moving their data and workloads into the cloud. The main reasons to move into the cloud are security, time-to-market, and a more flexible “pay-as-you-go” cost model which avoids overpaying for under-utilized machines. While all these reasons are compelling, the expectation is that a database runs in the cloud at least as well as (if not better) than on-premise. Specifically, customers expect a “database-as-a-service” to be highly available (e.g., 99.999% availability), support large databases (e.g., a 100TB OLTP database), and be highly performant. Furthermore, the service must be elastic and grow and shrink with the workload so that customers can take advantage of the pay-as-you-go model.

It turns out that meeting all these requirements is not possible in the cloud using the traditional monolithic database architecture. One issue is cost elasticity which never seemed to have been a consideration for on-premise database deployments. It can be prohibitively expensive to move a large database from one machine to another machine to support a higher or lower throughput and make the best use of the computing resources in a cluster. Another, more subtle issue is that there is a conflict between the goal to support large transactional databases and high availability: High availability requires a small mean-time-to-recovery which traditionally could only be achieved with a small database. This issue does not arise in on-premise database deployments because these deployments typically make use of special, expensive hardware for high availability (such as storage area networks or SANs); hardware which is not available in the cloud. Furthermore, on-premise deployments control the software update cycles and carefully plan downtimes; this planning is typically not possible in the cloud.

To address these challenges, there has been research on new OLTP database system architectures for the cloud over the last ten years; e.g., [5, 8, 16, 17]. One idea is to decompose the functionality of a database management system and deploy the compute services (e.g., transaction processing) and storage services (e.g., checkpointing and recovery) independently. The first commercial system that adopted this idea is Amazon Aurora [20].



Figure 2: Socrates Architecture

nodes cache data pages in main memory and on SSD in a resilient buffer pool extension (Section 3.3).

The second tier of the Socrates architecture is the XLOG service. This tier implements the “separation of log” principle, motivated in Section 4.1.4. While this separation has been proposed in the literature before [6, 8], this separation of log differentiates Socrates from other cloud database systems such as Amazon Aurora [20]. The XLOG service achieves low commit latencies and good scalability at the storage tier (scale-out). Since the Primary processes all updates (including DML operations), only the Primary writes to the log. This single writer approach guarantees low latency and high throughput when writing to the log. All other nodes (e.g., Secondaries) consume the log in an asynchronous way to keep their copies of data up to date.

The third tier is the storage tier. It is implemented by Page Servers. Each Page Server has a copy of a partition of the database, thereby deploying a scale-out storage architecture which, as we will see, helps to bound all operations as postulated in Section 4.1.2. Page Servers play two important roles: First, they serve pages to Compute nodes. Every Compute node can request pages from a Page Server, following a shared-disk architecture (Section 4.1.3). We are currently working on implementing bulk operations such as bulk loading, index creation, DB reorgs, deep page repair, and table scans in Page Servers to further offload Compute nodes as described in Section 4.1.5. In their second role, Page Servers checkpoint data pages and create backups in XStore (the

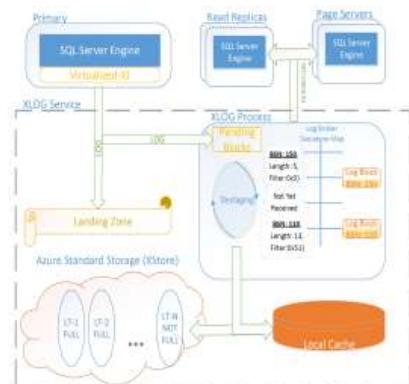


Figure 3: XLOG Service

fourth tier). Like Compute nodes, Page Servers keep all their data in main memory or locally attached SSDs for fast access.

The fourth tier is the Azure Storage Service (called XStore), the existing storage service provided by Azure independently of Socrates and SQL DB. XStore is a highly scalable, durable, and cheap storage service based on hard disks. Data access is remote and there are throughput and latency limits imposed by storage at that scale and price point. Separating Page Servers with locally attached, fast storage from durable, scalable, cheap storage implements the design principles outlined in Section 4.1.1.

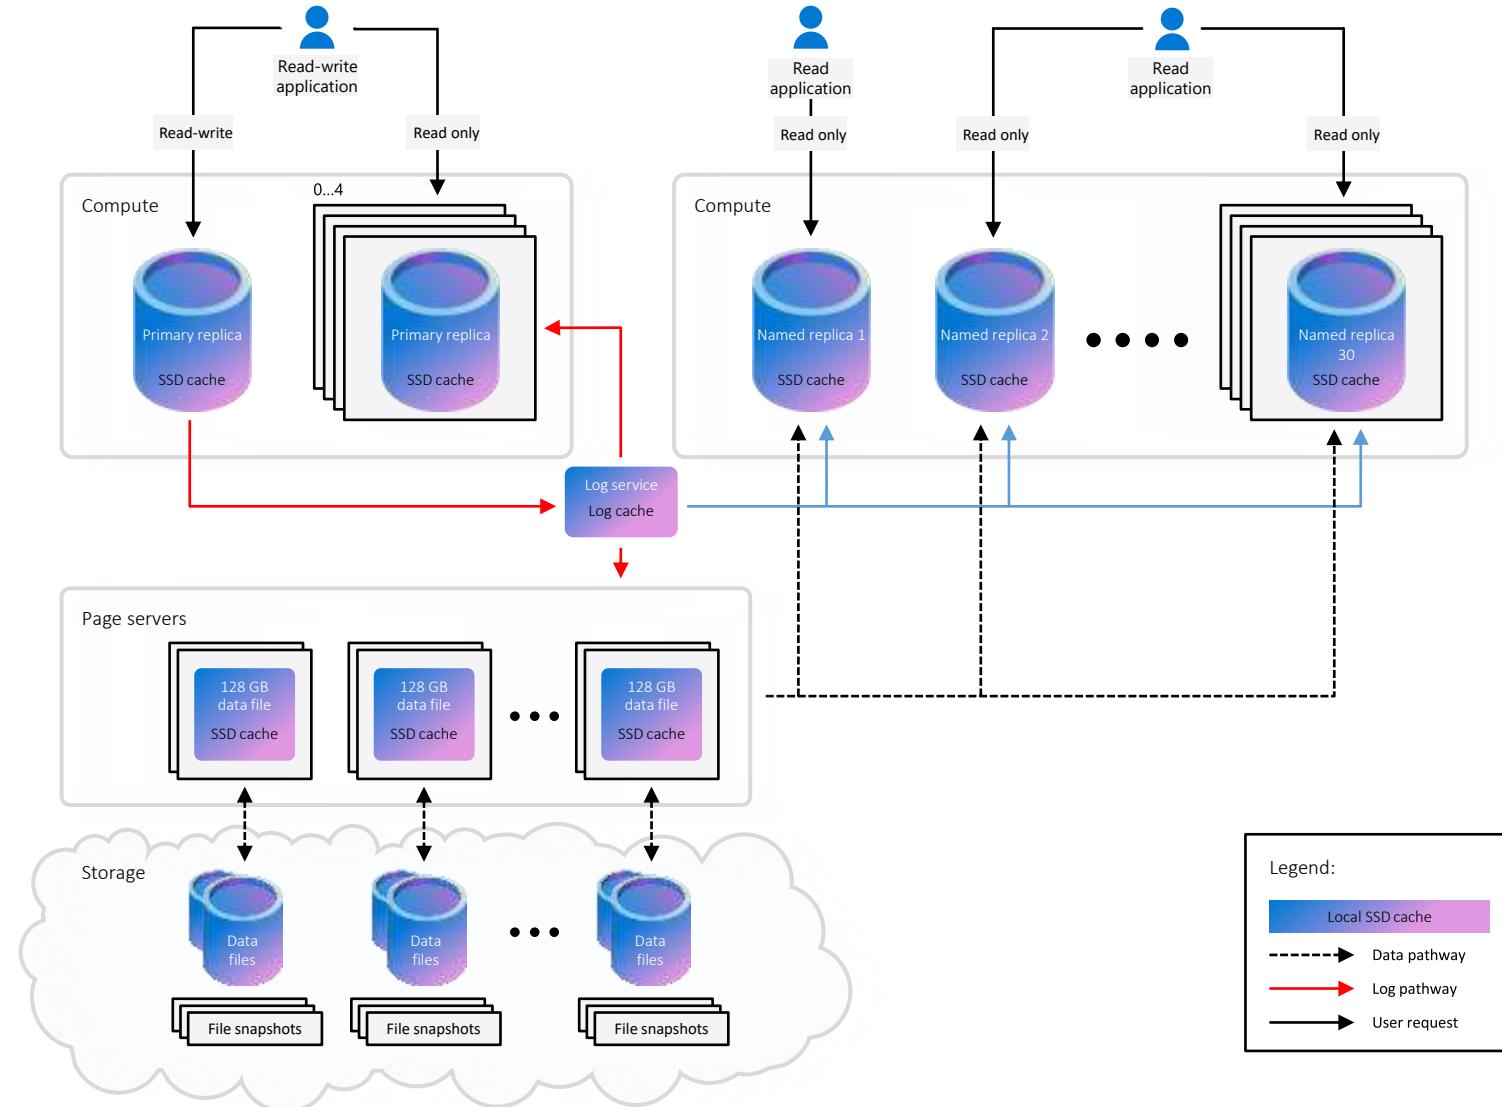
Compute nodes and Page Servers are stateless. They can fail at any time without data loss. The “truth” of the database is stored in XStore and XLOG. XStore is highly reliable and has been used by virtually all Azure customers for many years without data loss. Socrates leverages this robustness. XLOG is a new service that we built specifically for Socrates. It has high performance requirements, must be scalable, affordable, and must never lose any data. We describe our implementation of XLOG, Compute nodes, and Page Servers in more detail next.

4.3 XLOG Service

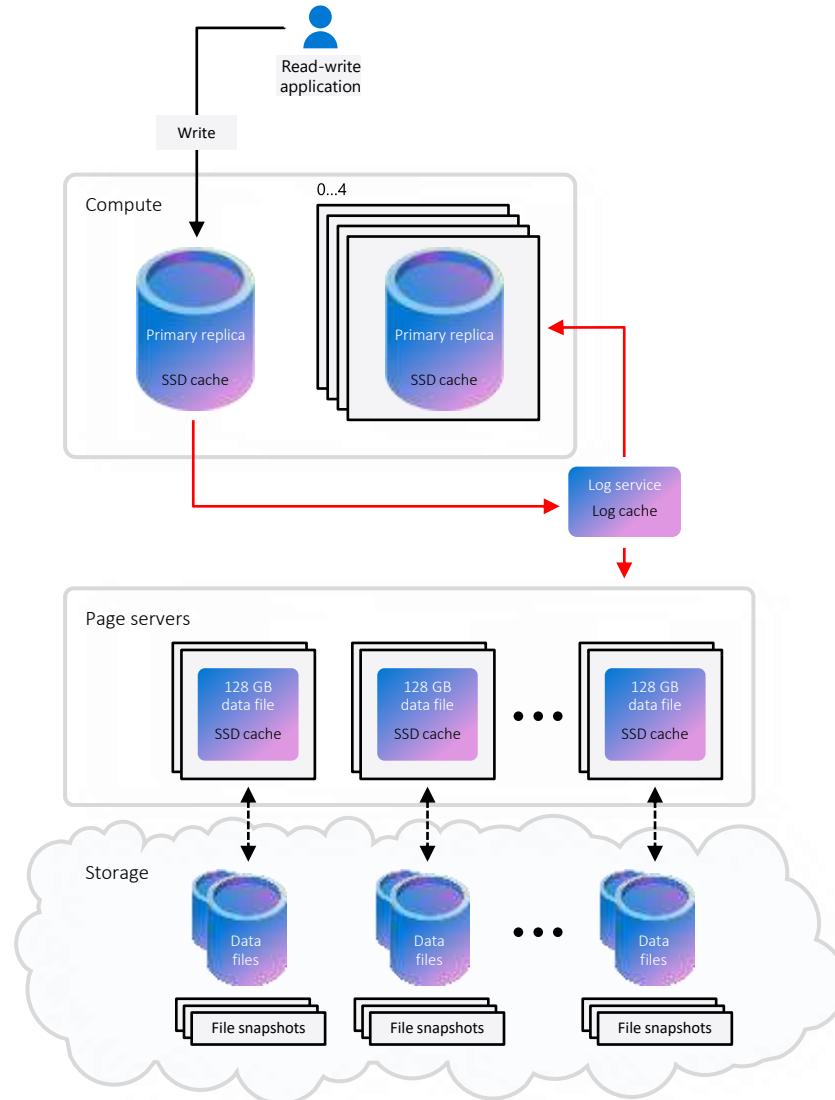
Figure 3 shows the internals of the XLOG Service. Starting in the upper left corner of Figure 3, the Primary Compute node writes log blocks directly to a landing zone (LZ) which is a fast and durable storage service that provides strong guarantees on data integrity, resilience and consistency; in other words, a storage service that has SAN-like capabilities.

Built-for-the cloud architecture enables near-limitless growth potential

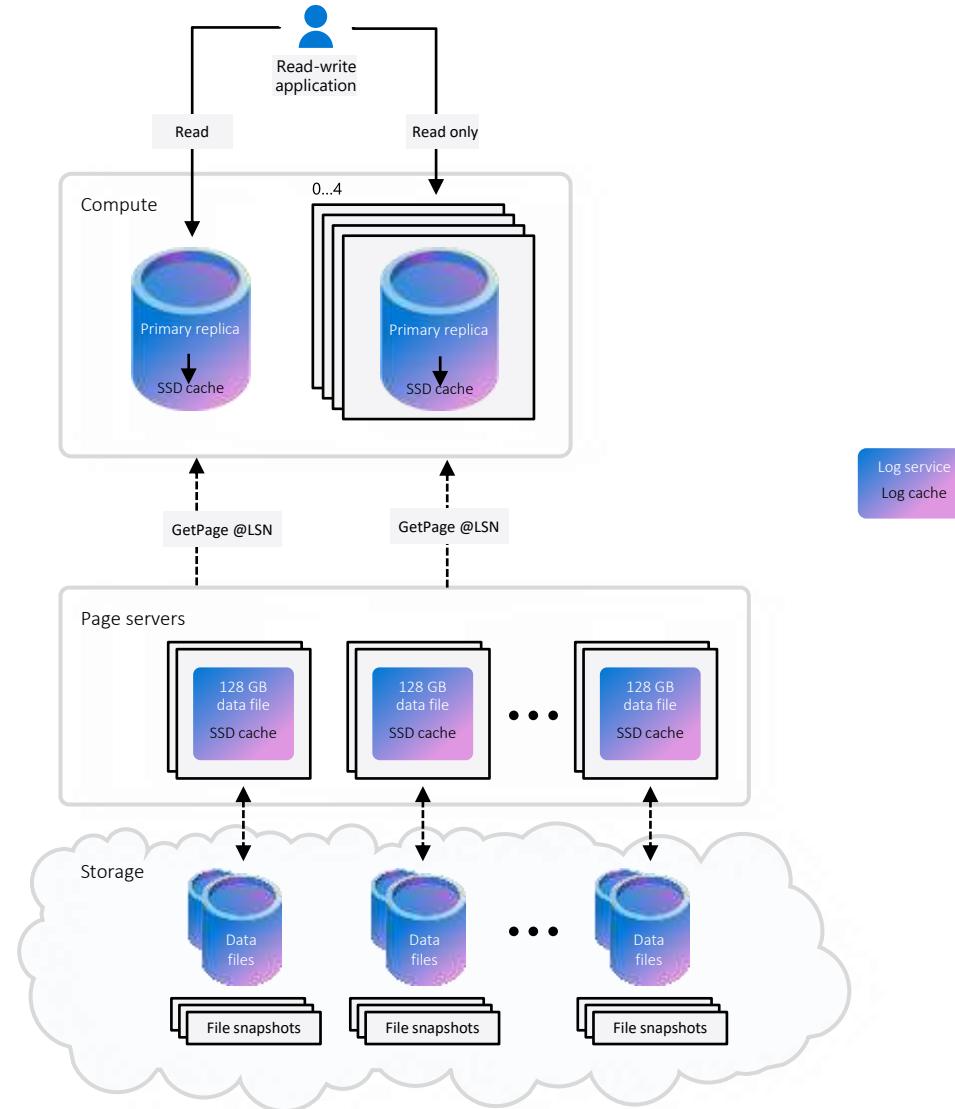
Distributed, scale-out and resilient



“Write path”



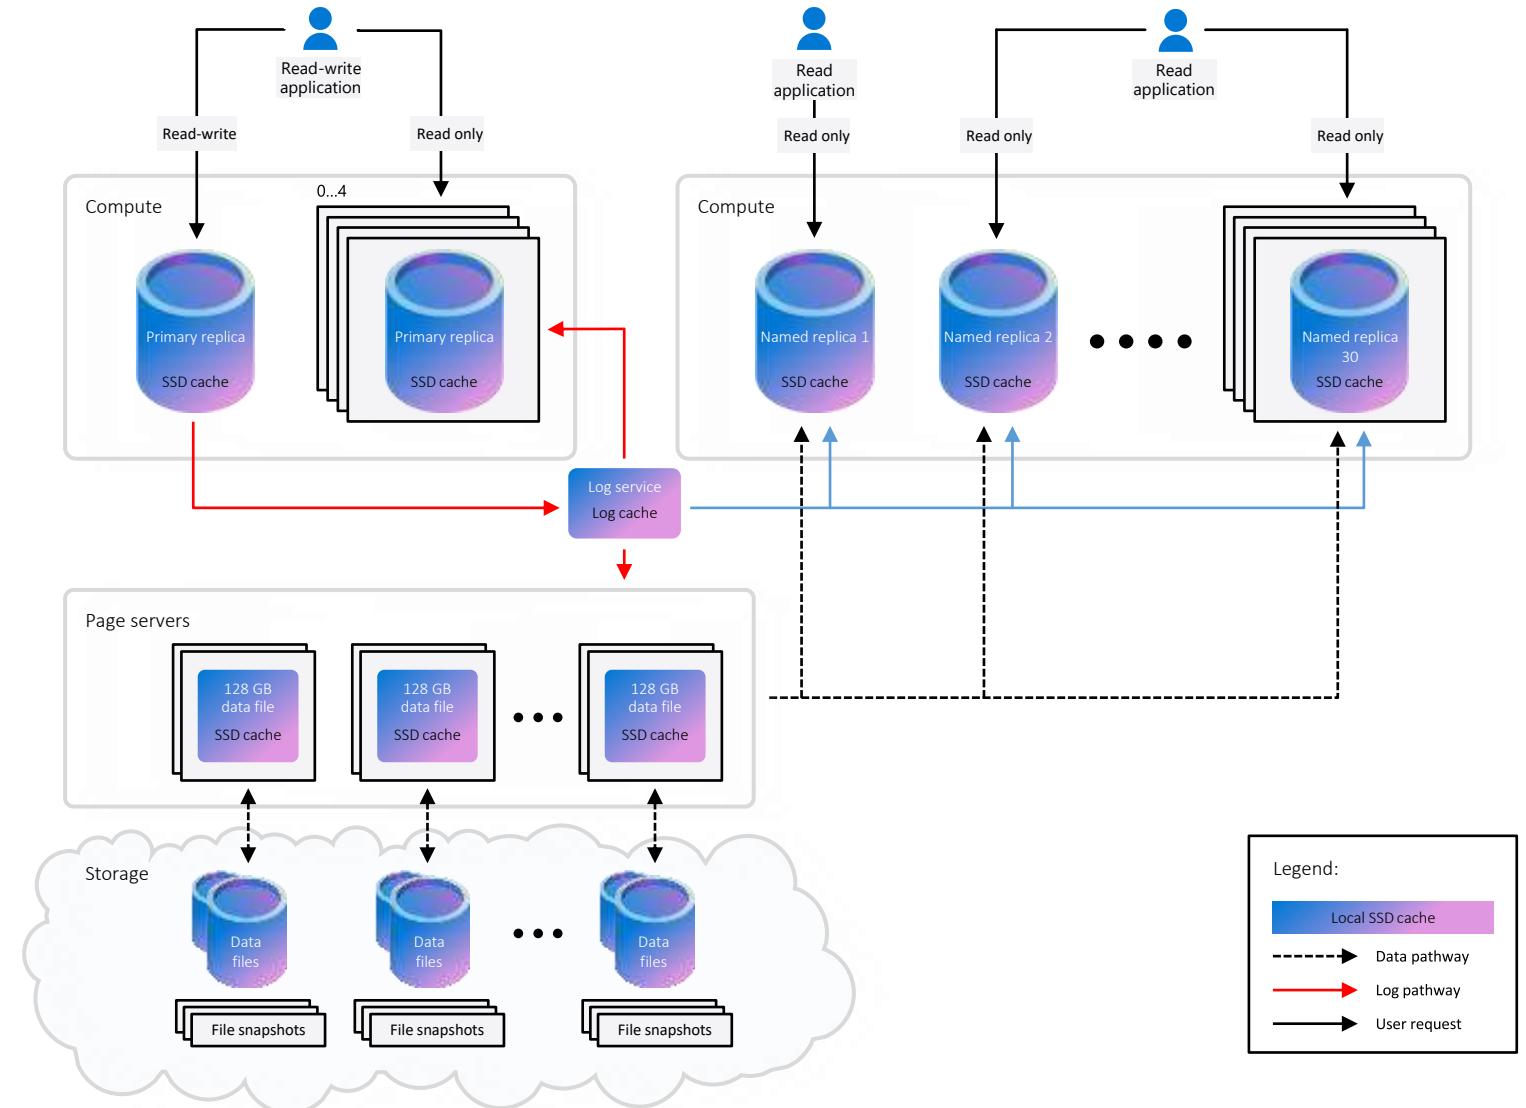
“Read path”



Built-for-the cloud architecture enables near-limitless growth potential

Distributed, scale-out and resilient

- Separation of compute & storage
- Smart SSD cache for compute
- Externalized log service
- Paired page servers, fully covering SSD cache
- Redundant data and log in Azure Storage
- 0 to 4 secondary HA replicas
- 0-30 named replicas for read scale
- Fast Backup/Restore via storage snapshots



Quiz time!



What is the maximum number of DB replicas (in a single server) which can be associated with the same underlying Hyperscale storage?



And the Answer
is

155



<https://aka.ms/hsignite2023>

Announcing new lower price tier for cloud native applications on Azure SQL



Asad Khan

November 14th, 2023 | 0 | 0



Today we are announcing a new pricing tier for Azure SQL Database targeted at developers building cloud native applications. It is based on Azure SQL Database Hyperscale tier, which offers awesome cloud scalability with one of the lowest TCO in the industry. Azure SQL Database Hyperscale delivers outstanding performance and scalability through adaptive, cloud native resource management, optimized for the workload patterns of each database application. Hyperscale storage scales automatically, providing consistent performance at all scales. Azure SQL Database Hyperscale provides the best database capabilities to build cloud native applications, including abilities to create APIs for your database using Data API builder, calling external endpoints, integrating with Azure functions, and more.

Market Leading Database Capabilities

Azure SQL Database Hyperscale brings three decades of SQL innovation to developers, along with the cloud scale and the latest AI capabilities. Some salient features that differentiate SQL Hyperscale database layer from the competition include:

- **Storage snapshot** capability in Hyperscale provides fast backup and restore, that meets the aggressive requirements of cloud native workloads. **Support for 30 Named Replicas allowing different compute sizes than the primary for each replica** is not offered in any other cloud database service. This is a huge advantage for highly transactional and read-scalable workloads both in terms of performance as well as cost.

<https://aka.ms/hspice2023>

Frequently Asked Questions

When does the change take effect, and what does it impact?

The pricing changes take effect on December 15th, 2023 at 00:00 hours UTC. The changes will apply to the following resources created on or after Dec 15th, 2023:

1. Any newly created Hyperscale single provisioned compute databases.
2. Any (existing or new) Hyperscale single [serverless](#) databases (currently in [preview](#)).
3. Any (existing or new) [Hyperscale elastic pools](#) (currently in [preview](#)).
4. Any newly created Hyperscale elastic pooled databases (currently in preview).

What happens to my existing Hyperscale resources?

To start with, nothing changes till December 15th, 2023. Here's what will happen starting December 15th, 2023:

1. To provide a seamless experience without any impact on existing workloads, all existing Hyperscale single databases with provisioned compute created before December 15th, 2023, will continue to be billed at the existing rates, for a period of up to 3 years (ending December 14th, 2026). Customers will receive further notice of the pricing change to their Hyperscale databases in advance of December 14th, 2026.
2. All existing Hyperscale single databases with serverless compute (currently in [preview](#)) will automatically switch to the new pricing starting December 15th, 2023.
3. All existing Hyperscale elastic pools (currently in [preview](#)) will automatically switch to the new pricing starting December 15th, 2023.

What if I update / change a database to a different deployment option?

Hyperscale allows [seamless, rapid scaling](#) of the database compute. You can also scale a Hyperscale database to move from provisioned compute to serverless compute (or the other way around). You can add an existing Hyperscale database to an elastic pool or move an elastic pooled database out of the elastic pool to a single database. Here's how your costs are impacted if you perform any of these changes on or after December 15th, 2023:

Change	Impact
Hyperscale (serverless single, or elastic pooled) database is changed to a Hyperscale single database with provisioned compute.	The final cost of the database will be based on when the database was created. If the database was created prior to December 15 th , 2023, it will be billed as per the existing pricing. If the database was created on or after December 15 th , 2023, it will be billed as per the new pricing.

Azure SQL Database Hyperscale outperforms

Build new, highly scalable cloud applications with the performance and security of Azure SQL *for the price of commercial open-source databases*

Principled Technologies [published](#) a study where they tested throughput performance between Azure SQL Database Hyperscale and Amazon Aurora PostgreSQL.

SQL Database Hyperscale emerged as the price-performance leader for mission-critical workloads, delivering up to 68 percent more performance per dollar than Amazon Aurora PostgreSQL.¹

[Blog: <https://devblogs.microsoft.com/azure-sql/build-highly-scalable-ai-ready-applications-on-azure-sql-database-hyperscale/>](https://devblogs.microsoft.com/azure-sql/build-highly-scalable-ai-ready-applications-on-azure-sql-database-hyperscale/)



Database performance per dollar comparison.
Higher is better.

Price-performance claims based on data from a study commissioned by Microsoft and conducted by Principled Technologies in December 2023. The study compared performance and price performance between a 16 vCore and 32 vCore Azure SQL Database using premium-series hardware on the Hyperscale service tier and the db.r6i.4xlarge and db.r6i.8xlarge offerings for Amazon Web Services Aurora PostgreSQL I/O-Optimized (AWS Aurora). Benchmark data is taken from a Principled Technologies report which used the HammerDB TPROC-C benchmark. The TPROC-C workload is derived from the TPC-C Benchmark and results were obtained with the HammerDB TPROC-C workload. The HammerDB TPROC-C workload is derived from the TPC-C benchmark and is not comparable to published TPC-C Benchmark results, as this implementation does not comply with all requirements of the TPC Benchmark. Price-performance is calculated by Principled Technologies as the cost of running the cloud platform continuously divided by new orders per minute throughput, based upon the standard. Prices are based on publicly available US pricing in East US 1 for Azure SQL Database and US East for AWS Aurora as of December 2023. Performance and price-performance results are based upon the configurations detailed in the Principled Technologies report. Actual results and prices may vary based on configuration and region.



Demo

Getting started with Hyperscale



mylogicalserver

SQL server



Search



+ Create database

+ New elastic pool

+ New dedicated SQL pool (formerly SQL DW)

Import database

Reset password



Essentials

JSON View

Notifications (0)

Features (6)

All

Security (4)

Performance (1)

Recovery (1)



Microsoft Entra admin

Allows you to centrally manage identity and access to your Azure SQL databases.

CONFIGURED



Microsoft Defender for SQL

Vulnerability Assessment and Advanced Threat Protection.

CONFIGURED



Automatic tuning

Monitors and tunes your database automatically to optimize performance.

CONFIGURED



Auditing

Track database events and writes them to an audit log in Azure storage.

NOT CONFIGURED

Available resources

Filter by name

All types





Demo

What's inside a Hyperscale DB?



File Edit View Query Project Tools Window Help

New Query MDX DMX XMLA DAX 1801

mynewDB1 Execute

Object Explorer

Connect mylogicalserver.database.windows.net

- Databases (filtered)
- System Databases
- myGeneralPurposeDB

Security

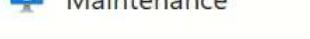
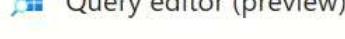
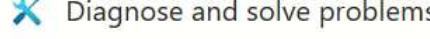
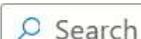
Integration Services Catalogs

SQLQuery9...om (141)* SQLQuery6...om (143)* SQLQuery5...om (140)*

```
SELECT name,
       physical_name,
       size * 8 / 1024.0 / 1024.0 AS SizeInGB,
       FILEPROPERTY(name, 'SpaceUsed') * 8 / 1024.0 as SpaceUsedMB,
       *
FROM sys.database_files
WHERE type_desc IN ('ROWS', 'LOG');
GO

SELECT *
FROM sys.dm_db_file_space_usage;
GO
```

142 %



Start working with your database

Connect to your database and start working with data with a few simple steps. [Learn more](#)



Configure access

Configure network access to your SQL server. [Learn more](#)

[Configure](#)

Connect to application

Use connection strings to connect to your SQL database from your applications and favorite tools.

[See connection strings](#)



Demo

Migrating existing DBs to Hyperscale





myGeneralPurposeDB (mylogicalserver/myGeneralPurposeDB) | Compute + storage





- [Overview](#)
- [Activity log](#)
- [Tags](#)
- [Diagnose and solve problems](#)
- [Query editor \(preview\)](#)

Settings

Compute + storage

Connection strings

Maintenance

Properties

Locks

Data management

Replicas

Sync to other databases

Integrations

Azure Synapse Link

Stream analytics (preview)

Add Azure AI Search



Service tier

General Purpose (Most budget friendly)

[Compare service tiers](#)

Compute tier

Provisioned - Compute resources are pre-allocated. Billed per hour based on vCores configured.

Serverless - Compute resources are auto-scaled. Billed per second based on vCores used.

Compute Hardware

Select the hardware configuration based on your workload requirements. Availability of compute optimized, memory optimized, and confidential computing hardware depends on the region, service tier, and compute tier.

Hardware Configuration

Standard-series (Gen5)

up to 128 vCores, up to 625 GB memory

[Change configuration](#)

Save money

Already have a SQL Server License? Save with a license you already own with Azure Hybrid Benefit. Actual savings may vary based on region and performance tier. [Learn more](#)

Yes

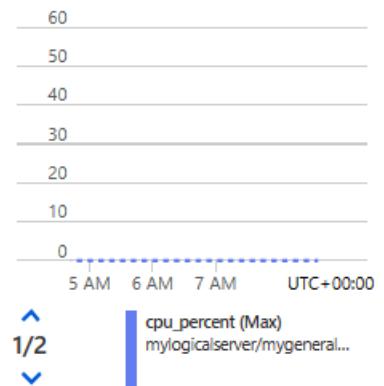
No

vCores [Compare vCore options](#)

24

Data max size (GB) ⓘ

4096



Cost summary

General Purpose (GP_Gen5_24)

Cost per vCore (in USD) **184.09**

vCores selected **x 24**

Cost per GB (in USD) **0.12**

Max storage selected (in GB) **x 5120**

ESTIMATED COST / MONTH

5007.05 usd

 Search

Feedback

Overview

Activity log

Tags

Diagnose and solve problems

Query editor (preview)

Resource visualizer

Settings

Compute + storage

Connection strings

Maintenance

Properties

Locks

Data management

Replicas

Sync to other databases

Integrations

Azure Synapse Link

Stream analytics (preview)

Service and compute tier

Select from the available tiers based on the needs of your workload. The vCore model provides a wide range of configuration controls and offers Hyperscale and Serverless to automatically scale your database based on your workload needs. Alternately, the DTU model provides set price/performance packages to choose from for easy configuration. [Learn more](#)

SQL Database Hyperscale: Low price, high scalability, and best feature set. [Learn more](#)

Service tier

General Purpose (Most budget friendly)

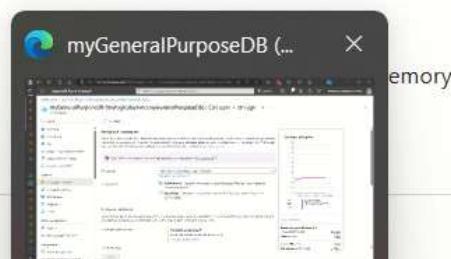
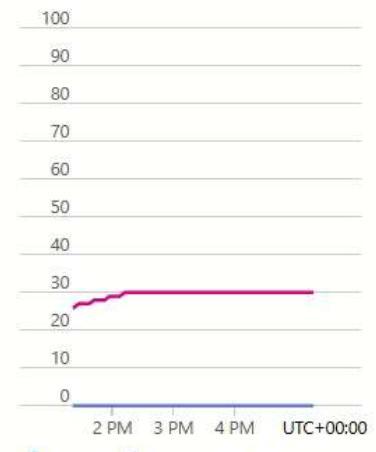
[Compare service tiers](#)

Provisioned - Compute resources are pre-allocated. Billed per hour based on vCores configured.

Serverless - Compute resources are auto-scaled. Billed per second based on vCores used.

Compute Hardware

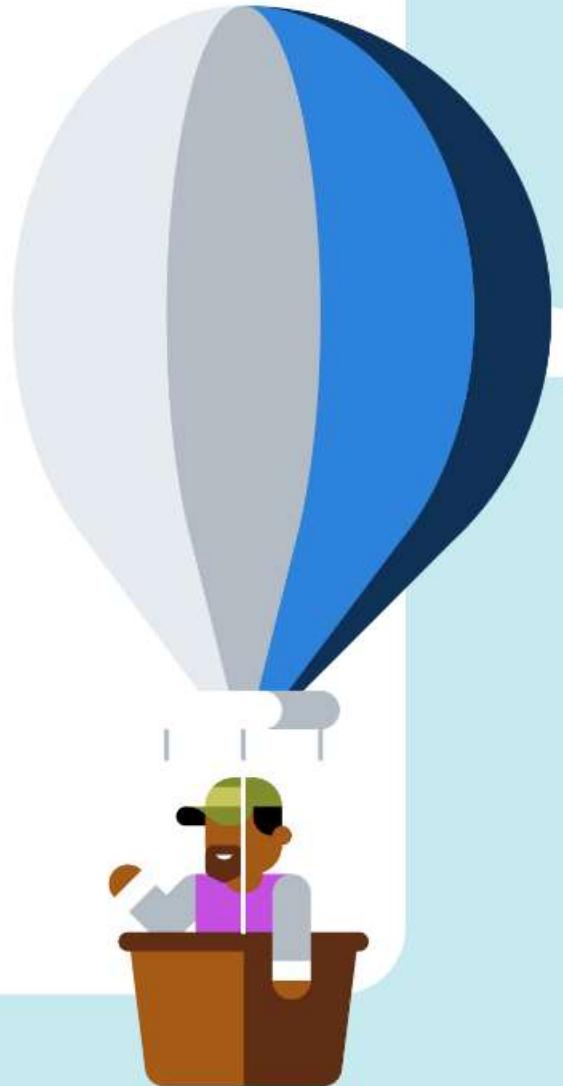
Select the hardware configuration based on your workload requirements. Availability of compute optimized, memory optimized, and confidential computing hardware depends on the region, service tier, and compute tier.

Hardware Configuration**Save money** Apply**Database utilization****Cost summary****General Purpose (GP_Gen5_24)**Cost per **vCore** (in USD)**184.09****vCores selected****x 24**Cost per **GB** (in USD)**0.12****Max storage selected (in GB)****x 5120**



Demo

Predictably quick scaling with Hyperscale





myNewDB1 (mylogicalserver/myNewDB1)

SQL database



Search



Copy

Restore

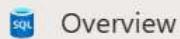
Export

Set server firewall

Delete

Connect with...

Feedback



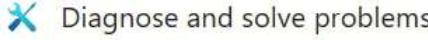
Overview



Activity log



Tags



Diagnose and solve problems



Query editor (preview)

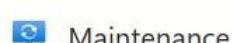
Settings



Compute + storage



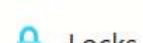
Connection strings



Maintenance



Properties



Locks

Data management



Essentials

[Getting started](#)[Monitoring](#)[Properties](#)[Features](#)[Notifications \(0\)](#)[Integrations](#)[Tutorials](#)

JSON View

Start working with your database

Connect to your database and start working with data with a few simple steps. [Learn more](#)



Configure access

Configure network access to your SQL server. [Learn more](#)

[Configure](#)

Connect to application

Use connection strings to connect to your SQL database from your applications and favorite tools.

[See connection strings](#)



myNewDB1 (mylogicalserver/myNewDB1)

SQL database



Search



Feedback



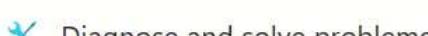
Overview



Activity log



Tags

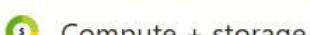


Diagnose and solve problems



Query editor (preview)

Settings



Compute + storage



Connection strings



Maintenance

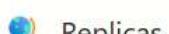


Properties



Locks

Data management



Replicas

Scaling from HS_Gen5_10 to HS_PRMS_8 database level [See more](#)

Essentials

[Getting started](#)[Monitoring](#)[Properties](#)[Features](#)[Notifications \(1\)](#)[Integrations](#)[Tutorials](#)[JSON View](#)

Start working with your database

Connect to your database and start working with data with a few simple steps. [Learn more](#)



Configure access

Configure network access to your SQL server. [Learn more](#)

[Configure](#)

Connect to application

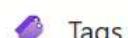
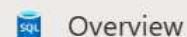
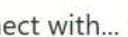
Use connection strings to connect to your SQL database from your applications and favorite tools.

[See connection strings](#)

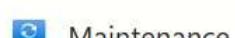


myNewDB1 (mylogicalserver/myNewDB1)

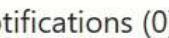
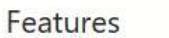
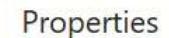
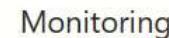
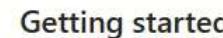
SQL database



Settings



Data management



JSON View

Scaling operation completed



Start working with your database

Connect to your database and start working with data with a few simple steps. [Learn more](#)



Configure access

Configure network access to your SQL server. [Learn more](#)

[Configure](#)

Connect to application

Use connection strings to connect to your SQL database from your applications and favorite tools.

[See connection strings](#)



Object Explorer

Connect

mylogicalserver.database.windows.net

Databases (filtered)

mylogicalserver.database.windows.net (SQL Server 12.0.2000.8 - arvindsh@microsoft.com)

System Databases

myNewDB1

Database Diagrams

Tables

Views

External Resources

Synonyms

Programmability

Query Store

Extended Events

Storage

Security

Security

Integration Services Catalogs

SQLQuery2...com (55)*

```
SELECT *  
FROM sys.dm_operation_status;
```

GO

SQLQuery1...om (157)*

```
SELECT DATABASEPROPERTYEX(db_name(), 'ServiceObjective');
```

GO

mylogicalserver.database.windows.net (SQL Server 12.0.2000.8 - arvindsh@microsoft.com)

142 %

Demo

Published characteristics (“limits”) of a Hyperscale DB

Hyperscale - provisioned compute - premium-series

Hyperscale premium-series (part 1 of 3)

Compute sizes (service level objectives, or SLOs) for Hyperscale premium-series databases follow the naming convention `HS_PRMS_` followed by the number of vCores.

The following table covers these SLOs: `HS_PRMS_2`, `HS_PRMS_4`, `HS_PRMS_6`, `HS_PRMS_8`, and `HS_PRMS_10`:

[Expand table](#)

vCores	2	4	6	8	10
Hardware	Premium-series	Premium-series	Premium-series	Premium-series	Premium-series
Memory (GB)	10.4	20.8	31.1	41.5	51.9
Columnstore support	Yes	Yes	Yes	Yes	Yes
In-memory OLTP storage (GB)	N/A	N/A	N/A	N/A	N/A
Max data size (TB)	100	100	100	100	100
Max log size (TB)	Unlimited	Unlimited	Unlimited	Unlimited	Unlimited
Tempdb max data size (GB)	64	128	192	256	320
Max local SSD IOPS ¹	8,000	16,000	24,000	32,000	40,000
Max log rate (MBps)	100	100	100	100	100
Local read IO latency ²	1-2 ms				
Remote read IO latency ²	1-4 ms				
Write IO latency ²	1-4 ms				
Storage type	Multi-tiered ³				

[https://aka.ms/
hslimits](https://aka.ms/hslimits)

Resource management in Azure SQL Database

Article • 12/28/2023 • 16 contributors

Feedback

In this article

Logical server limits

What happens when resource limits are reached

Resource consumption by user workloads and internal processes

Resource governance

Show 3 more

[https://aka.ms/
sqldbrg](https://aka.ms/sqldbrg)

Applies to:  Azure SQL Database

Azure SQL Database logical server ▾

This article provides an overview of resource management in Azure SQL Database. It provides information on what happens when resource limits are reached, and describes resource governance mechanisms that are used to enforce these limits.

For specific resource limits per pricing tier (also known as service objective) for single databases, refer to either DTU-based single database resource limits or vCore-based single database resource limits. For elastic pool resource limits, refer to either DTU-based elastic pool resource limits or vCore-based elastic pool resource limits. For Azure Synapse Analytics dedicated SQL pool limits, see [capacity limits](#) and [memory and concurrency limits](#).

File Edit View Query Project Tools Window Help

New Query MDX DMX XMLA DAX

mynewDB1 Execute

Object Explorer

Connect mylogicalserver.database.windows.net

- Databases (filtered)
- System Databases
- myNewDB1
 - Database Diagrams
 - Tables
 - Views
 - External Resources
 - Synonyms
 - Programmability
 - Query Store
 - Extended Events
 - Storage
 - Security
- Security
- Integration Services Catalogs

SQLQuery2..om (147)*

```
SELECT db_name(),
       database_id,
       DATABASEPROPERTYEX(db_name(), 'ServiceObjective'),
       slo_name,
       initial_db_file_size_in_mb / 1024.0 AS initial_db_file_size_in_
       db_file_growth_in_mb / 1024.0 as db_file_growth_in_gb,
       primary_group_max_workers,
       primary_group_max_io,
       *
FROM sys.dm_user_db_resource_governance;
```



Object Explorer

Connect

- mylogicalserver.database.windows.net
 - Databases (filtered)
 - System Databases
 - myNewDB1
 - Database Diagrams
 - Tables
 - Views
 - External Resources
 - Synonyms
 - Programmability
 - Query Store
 - Extended Events
 - Storage
 - Security
 - Security
 - Integration Services Catalogs

SQLQuery2...om (147)*

```
SELECT process_memory_limit_mb / 1024.0 as process_memory_limit_gb  
FROM sys.dm_os_job_object;
```

Hyperscale FAQs

Why decouple compute and storage?

Why this complex architecture?

Why is the log generation rate for Hyperscale limited to 100 MB / sec?

Why can't I restore a backup of a Hyperscale DB elsewhere?

What does a PAGEIOLATCH_* wait mean in a Hyperscale database?

Does adding more vCores in Hyperscale always lead to better performance?

Let's talk performance tuning

Specifics of workload

The screenshot shows the TPROC-C workload generator interface. On the left, there's a list of supported databases: Benchmark, SQL Server, TPROC-C, Oracle, DB2, MySQL, PostgreSQL, and MariaDB. The 'Benchmark' tab is selected. In the center, there's a script editor window displaying a TPC-H benchmark script. A configuration dialog box is open over the script, titled 'Virtual User Options'. It contains fields for 'Virtual Users' (set to 250), 'User Delay(ms)', 'Repeat Delay(ms)', and 'Iterations'. Below these are several checkboxes: 'Show Output', 'Log Output to Temp', 'Use Unique Log Name', 'No Log Buffer', and 'Log Timestamps'. At the bottom of the dialog are 'OK' and 'Cancel' buttons. To the right of the script editor, another configuration dialog is open, titled 'Microsoft SQL Server TPROC-C Driver Options'. It includes fields for 'SQL Server' (set to 'atabase.windows.net'), 'TCP' (set to 1433), 'Azure' (checkbox checked), 'Encrypt Connection' (checkbox checked), 'Trust Server Certificate' (checkbox checked), 'SQL Server ODBC Driver' (set to 'ODBC Driver 18 for SQL Server'), 'Authentication' (radio button selected for 'Windows Authentication'), 'SQL Server User ID' (empty), 'SQL Server User Password' (empty), 'TPROC-C SQL Server Database' (set to 'TPCCHyperscale'), 'TPROC-C Driver Script' (radio button selected for 'Timed Driver Script'), 'Total Transactions per User' (set to 10000000), 'Exit on SQL Server Error' (checkbox checked), 'Keying and Thinking Time' (checkbox checked), 'Checkpoint when complete' (checkbox checked), 'Minutes of Rampup Time' (set to 1), 'Minutes for Test Duration' (set to 30), 'Use All Warehouses' (checkbox checked), 'Time Profile' (checkbox checked), 'Asynchronous Scaling' (checkbox checked), 'Asynch Clients per Virtual User' (set to 10), 'Asynch Client Login Delay' (set to 1000), 'Asynchronous Verbosity' (checkbox checked), and 'XML Connect Pool' (checkbox checked). At the bottom of this dialog are 'OK' and 'Cancel' buttons.

```
1 #!/usr/local/bin/tclsh8.6
2 #EDITABLE OPTIONS#####
3 set library tdbcodbc # SQL Server Library
4 set version 1.1.1 # SQL Server Library Version
5 set total_iterations 10000000 # Number of transactions before logging off
6 set RAISEERROR "false" # Exit script on SQL Server error (true or false)
7 set KEYANDTHINK "false" # Time for user thinking and keying (true or false)
8 set CHECKPOINT "false" # Perform SQL Server checkpoint when complete (true or false)
9 set rampup 1; # Rampup time in minutes before first Transaction Count is taken
10 set duration 30; # Duration in minutes before second Transaction Count is taken
11 set mode "local" # HammerDB operational mode
12 set authentication "sql" # Authentication Mode (WINDOWS or SQL)

# Virtual User Options
Virtual Users: 250
User Delay(ms): 5
Repeat Delay(ms): 5
Iterations: 1
Show Output
Log Output to Temp
Use Unique Log Name
No Log Buffer
Log Timestamps
OK Cancel

31 if [ | chk_thread | eq "FALSE" ] |
32 error "SQL Server Timed Script must be run in Thread Enabled Interpreter"
33 |
34
35 proc connect_string | server port odbc_driver authentication uid pwd tcp azure db encrypt trust_cert |
36 # | $tpc eq "true" | set server tcp:$server:$port |
37 # | $string toupper $authentication | eq "WINDOWS" |
38 | set connection "DRIVER=$odbc_driver;SERVER=$server;TRUSTED_CONNECTION=YES"
39 | else |
40 | | $string toupper $authentication | eq "SQL" |
41 | | set connection "DRIVER=$odbc_driver;SERVER=$server;UID=$uid;PWD=$pwd"
42 | | else |
43 | | puts stderr "Error: neither WINDOWS or SQL Authentication has been specified"
44 | | set connection "DRIVER=$odbc_driver;SERVER=$server"
45 |
46
47 # | $azure eq "true" | | append connection ";" "DATABASE=$db"
48 # | $encrypt eq "true" | | append connection ";" "ENCRYPT=yes" | else | append connection ";" "ENCRYPT=no"
49 # | $trust_cert eq "true" | | append connection ";" "TRUSTSERVERCERTIFICATE=yes"
50 return $connection
```

Specifics

SQL DB HS - vCore -16

Client vCore- 16

DB size- 3.3 TB



Demo

sys.dm_db_resource_stats



```
DECLARE @db_max_log_rate AS FLOAT = (SELECT primary_max_log_rate
                                       FROM sys.dm_user_db_resource_governance
                                       WHERE database_id = db_id());

SELECT TOP 1 used_storage_mb,
       used_storage_mb / 1024.0 / 1024.0 AS UsedSpaceTiB,
       allocated_storage_mb,
       allocated_storage_mb / 1024.0 / 1024.0 AS AllocatedSpaceTiB,
       avg_data_io_percent,
       avg_log_write_percent,
       avg_log_write_percent * @db_max_log_rate / (100.0 * 1024.0 * 1024.0) AS LogWriteMiBPerSec,
       avg_cpu_percent AS avg_cpu_percent
FROM sys.dm_db_resource_stats
ORDER BY end_time DESC;
```

I

PAGEIOLATCH_*

PAGEIOLATCH_EX

Occurs when a task is waiting on a latch for a buffer that is in an I/O request. The latch request is in Exclusive mode - a mode used when the buffer is being written to disk. Long waits might indicate problems with the disk subsystem.

For more information, see [Slow I/O - SQL Server and disk I/O performance ↗](#).

```
SQLQuery3...ram (240)* ✘ × SQLQuery2...ram (417)* SQLQuery1...ram (177)*  
SELECT      wait_type,  
            wait_duration_ms,  
            resource_description  
FROM        sys.dm_os_waiting_tasks AS WT  
INNER JOIN  
            sys.dm_exec_sessions AS S  
ON        WT.session_id = S.session_id  
WHERE       S.is_user_process = 1  
  
go  
  
select * from sys.databases where database_id=  
go  
sp_helpfile
```

PAGEIOLATCH_* in context of Hyperscale

- Same definition as in SQL Server.
- Indicates remote IO to the Page Servers
- Occurs because of reasons such as:
 - BPool + RBPEX smaller than what the workload ideally needs.
 - Higher transaction rate which is slowing down Redo in Pageservers.

WRITELOG

WRITELOG

Occurs while waiting for a log flush to complete. Common operations that cause log flushes are transaction commits and checkpoints. Common reasons for long waits on WRITELOG are: disk latency (where transaction log files reside), the inability for I/O to keep up with transactions, or, a large number of transaction log operations and flushes (commits, rollback)

SQLQuery3...ram (240)* SQLQuery2...ram (417)* SQLQuery1...ram (177)*

```
SELECT      wait_type,
            wait_duration_ms,
            resource_description
FROM        sys.dm_os_waiting_tasks AS WT
INNER JOIN
            sys.dm_exec_sessions AS S
ON        WT.session_id = S.session_id
WHERE       S.is_user_process = 1
and wait_type like '%write%' I
go
```

WRITELOG in context of Hyperscale

- Same definition as in SQL Server.
- Here the transaction log that is written to is the “XLOG service”.
- This occurs in Hyperscale only if
 - Higher commits on the application
 - There are any bottlenecks with the Log Service:
 - This should be only for exceptional scenarios.
 - Usually, we'll know well before you know if there is such a bottleneck with the Log Service.
 - And if that's the case, we (Microsoft) will internally resolve it.

Quiz time!

Would you see a LOG_RATE_GOVERNOR wait type in the on-prem world?

SQLQuery4....ecuting...* SQLQuery3...ram (240)* SQLQuery2...ram (417)* SQLQuery1...ram (177)*

```
alter index PK_Stock on stock rebuild
```

I

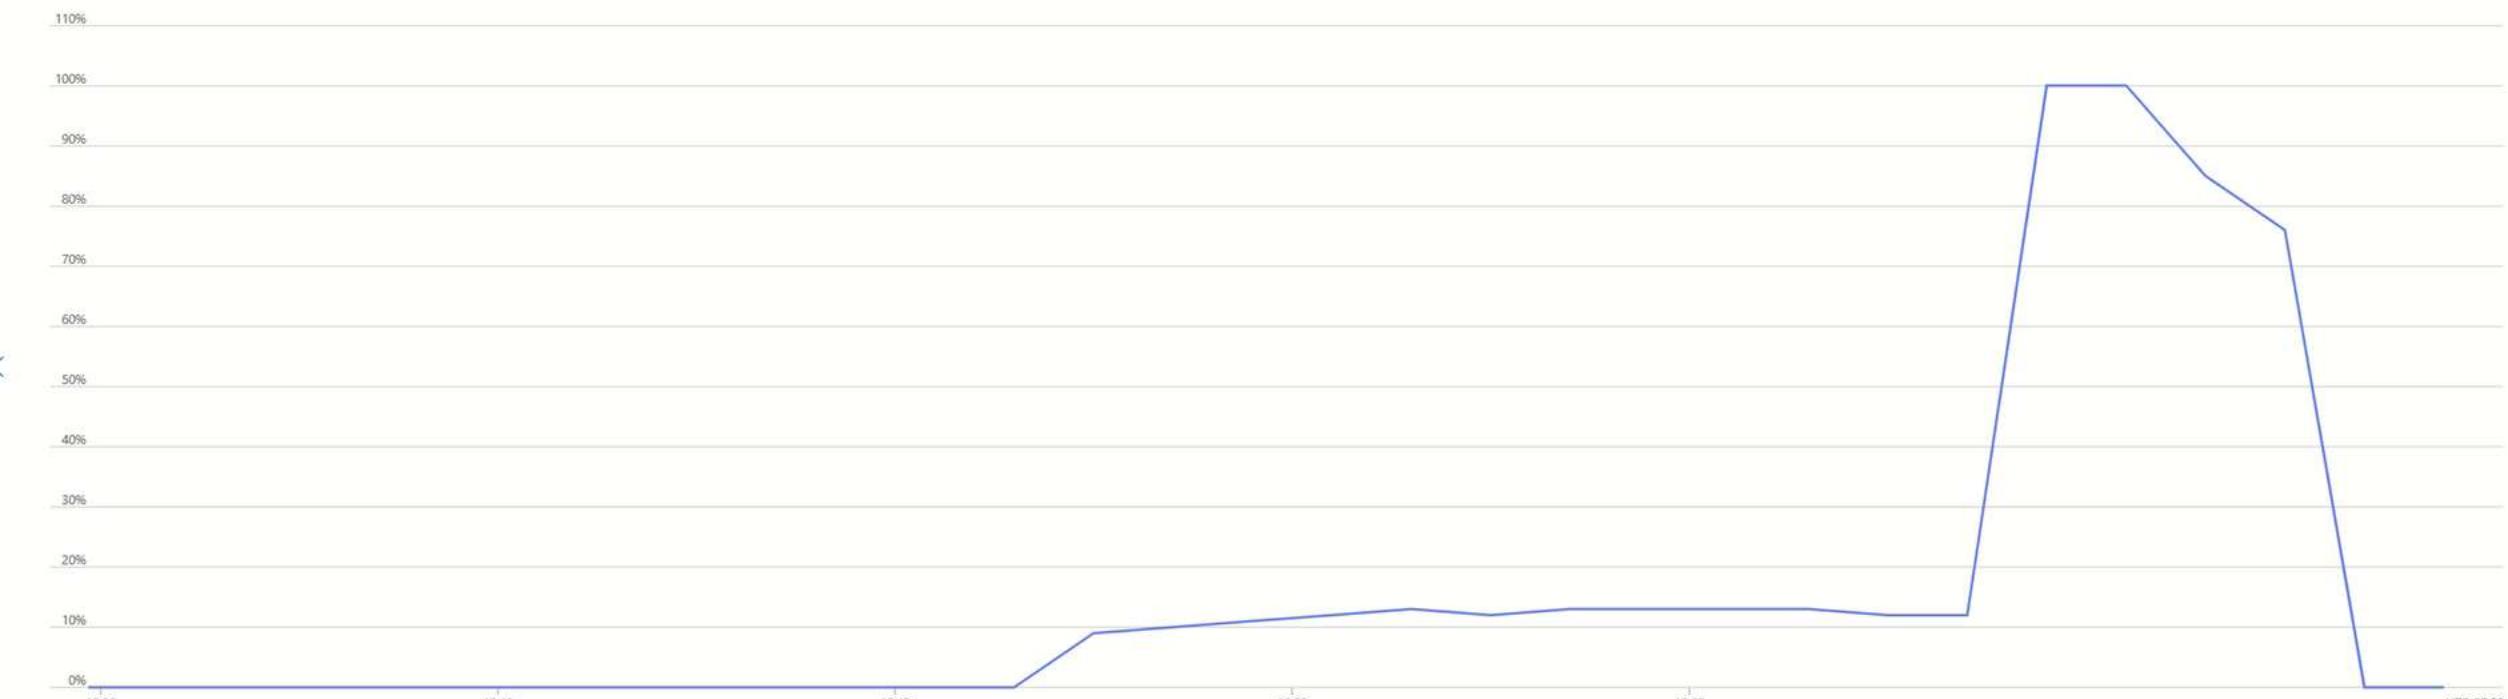
150 % <

Results Messages

Executing query... adbadram.database.windows.n... adbadram (810) TpcchHyperscale 00:01:43 | 0 rows

[New chart](#)  Refresh  Share  Feedback

Local Time: Last 30 minutes (1 minute)

Compute utilization  Add metric  Add filter  Apply splitting Line chart  Drill into Logs  New alert rule  Save to dashboard ... adbadram/TpccHyperscale, Log IO percentage, Max  Metric=Log IO percentage (Max), Resource=TpccHyperscale 100%

Azure SQL DB Hyperscale serverless

GA !!



Automatic scaling of compute and memory; log throughput independent of compute



Auto-scaling independence of the primary replica, high availability replicas, and named replicas



Automatic scaling of database storage up to 100 TB



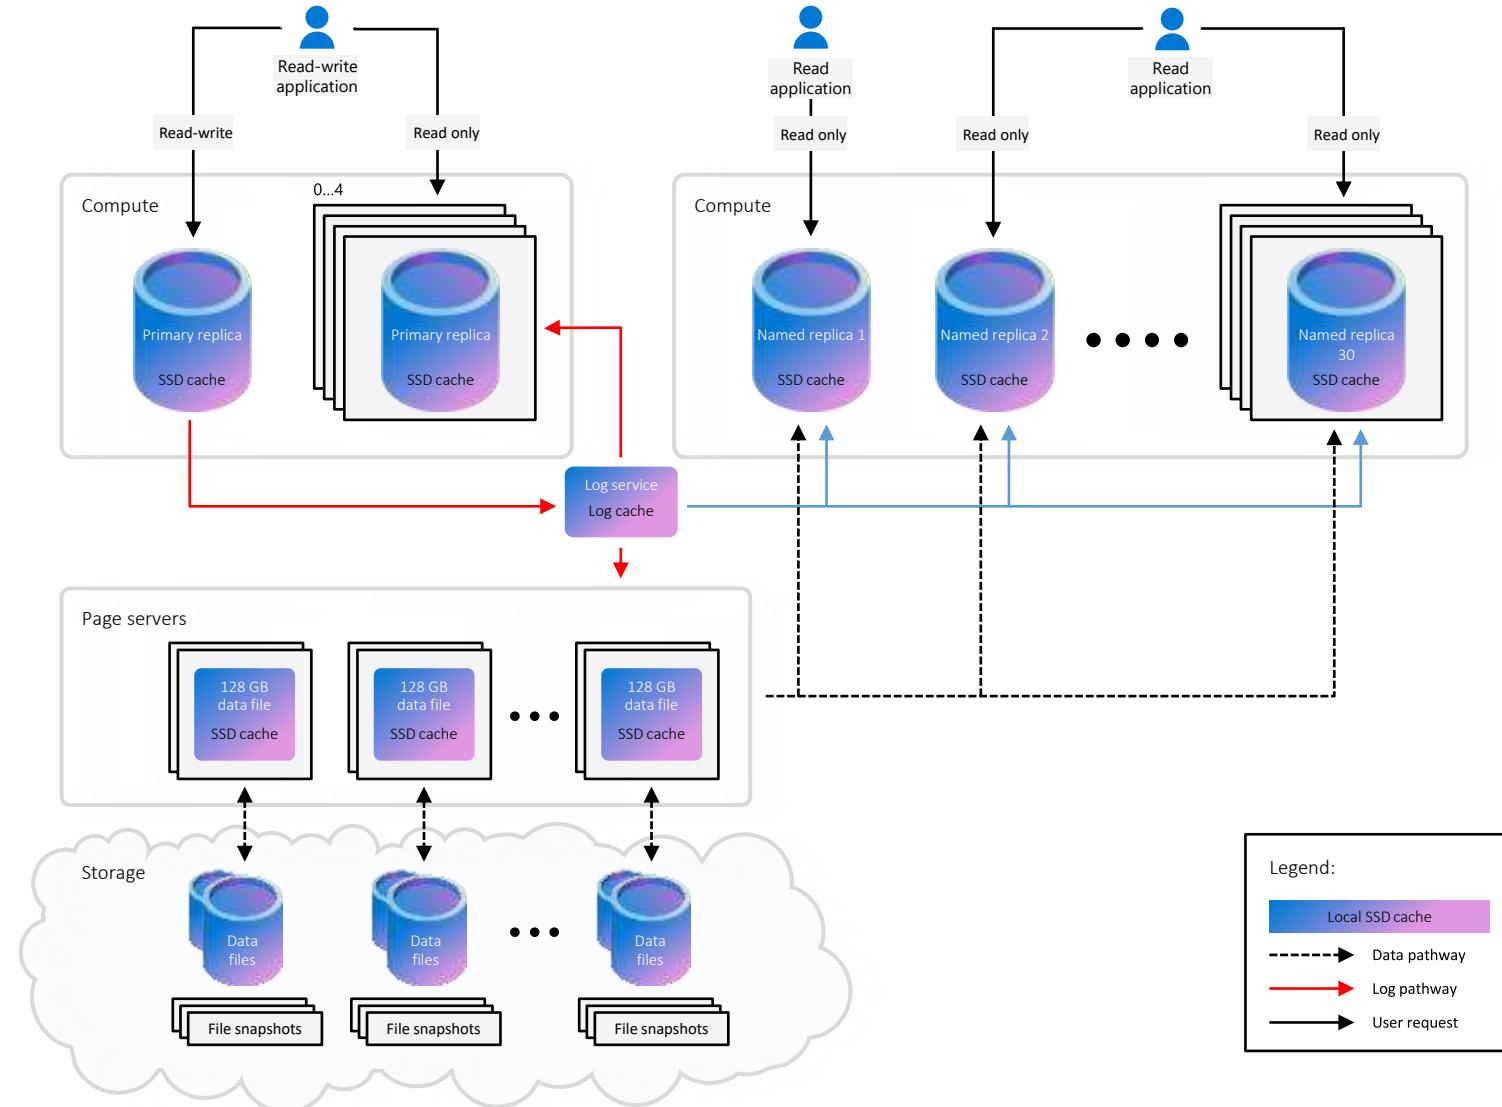
Auto-scaling independence of CPU and memory to match workload demand



Per-Second Billing of used resources

Built-for-the cloud architecture enables near-limitless growth potential

Distributed, scale-out and resilient





Demo

Virtually infinite read scale with Hyperscale + serverless



 Search[Copy](#) [Restore](#) [Export](#) [Set server firewall](#) [Delete](#) [Connect with...](#) [Feedback](#)[Overview](#)[Activity log](#)[Tags](#)[Diagnose and solve problems](#)[Query editor \(preview\)](#)

Settings

[Compute + storage](#)[Connection strings](#)[Maintenance](#)[Properties](#)[Locks](#)

Data management

[Replicas](#)

Essentials

[JSON View](#)[Getting started](#)[Monitoring](#)[Properties](#)[Features](#)[Notifications \(0\)](#)[Integrations](#)[Tutorials](#)

Start working with your database

Connect to your database and start working with data with a few simple steps. [Learn more](#)



Configure access

Configure network access to your SQL server. [Learn more](#)

[Configure](#)

Connect to application

Use connection strings to connect to your SQL database from your applications and favorite tools.

[See connection strings](#)



myNewDB1 (mylogicalserver/myNewDB1) | Replicas



Search



Create replica

Refresh

Feedback



Overview

Activity log

Tags

Diagnose and solve problems

Query editor (preview)

Settings

Compute + storage

Connection strings

Maintenance

Properties

Locks

Data management

Replicas

Integrations

Azure Synapse Link

Stream analytics (preview)

Add Azure AI Search

Power Platform

Want to enable disaster recovery at lower price? Create standby replica for General Purpose or Business Critical service tier, with savings of up to 40%. [Learn more](#)

Geo and named replicas for your database are listed below. Named replicas reside in the same region as the primary and enable load-balancing of read-only workloads and other scenarios. Geo replicas reside on a different logical server from the primary and protect against regional failures or prolonged data center outage. [Learn more](#)

Name ↑↓	Server ↑↓	Region ↑↓	HA replica count ↑↓	Failover poli... ↑↓	Pricing tier ↑↓	Replica state ↑↓	
myNewDB1	mylogicalserver	East US	1	None	Hyperscale: Premium-series, 8 vCores	Online	...
myNewDB1_NR1	mylogicalserver	East US	4		Hyperscale - Serverless: Standard-series (Gen5), ...	Readable	...
myNewDB1_NR10	mylogicalserver	East US	4		Hyperscale - Serverless: Standard-series (Gen5), ...	Readable	...
myNewDB1_NR11	mylogicalserver	East US	4		Hyperscale - Serverless: Standard-series (Gen5), ...	Readable	...
myNewDB1_NR12	mylogicalserver	East US	4		Hyperscale - Serverless: Standard-series (Gen5), ...	Readable	...
myNewDB1_NR13	mylogicalserver	East US	4		Hyperscale - Serverless: Standard-series (Gen5), ...	Readable	...
myNewDB1_NR14	mylogicalserver	East US	4		Hyperscale - Serverless: Standard-series (Gen5), ...	Readable	...
myNewDB1_NR15	mylogicalserver	East US	4		Hyperscale - Serverless: Standard-series (Gen5), ...	Readable	...
myNewDB1_NR16	mylogicalserver	East US	4		Hyperscale - Serverless: Standard-series (Gen5), ...	Readable	...
myNewDB1_NR17	mylogicalserver	East US	4		Hyperscale - Serverless: Standard-series (Gen5), ...	Readable	...
myNewDB1_NR18	mylogicalserver	East US	4		Hyperscale - Serverless: Standard-series (Gen5), ...	Readable	...

Behind the scenes...

```
(1..30) | % { $jobs += Start-Job -ScriptBlock
{ New-AzSqlDatabaseSecondary -ResourceGroupName
"myResourceGroup" -ServerName "mylogicalserver" -
DatabaseName "myNewDB1" -PartnerResourceGroupName
"myResourceGroup" -PartnerServerName
"mylogicalserver" -PartnerDatabaseName
("myNewDB1_NR" + $args[0]) -SecondaryType Named -
SecondaryServiceObjectiveName HS_S_Gen5_2 -
MaxVcore 2 -MinVcore 0.5 -
HighAvailabilityReplicaCount 4
} -ArgumentList ($_) }
```

 **myNewDB1_NR20** (mylogicalserver/myNewDB1_NR20) ⚡ ⭐ ...

SQL database

Search

Copy Restore Export Set server firewall Delete Connect with... Feedback

Overview Essentials JSON View

Activity log Tags Diagnose and solve problems Query editor (preview)

Getting started Monitoring Properties Features Notifications (0) Integrations Tutorials

Start working with your database

Connect to your database and start working with data with a few simple steps. [Learn more](#)

 **Configure access**
Configure network access to your SQL server. [Learn more](#)

Configure

 **Connect to application**
Use connection strings to connect to your SQL database from your applications and favorite tools.

See connection strings

 **Start developing**
Work in your database by using tools to add, modify and query data. [Compare tools](#)

Open Azure Data Studio

Open in Visual Studio

Open in Visual Studio Code

Settings

Compute + storage Connection strings Maintenance Properties Locks

Data management

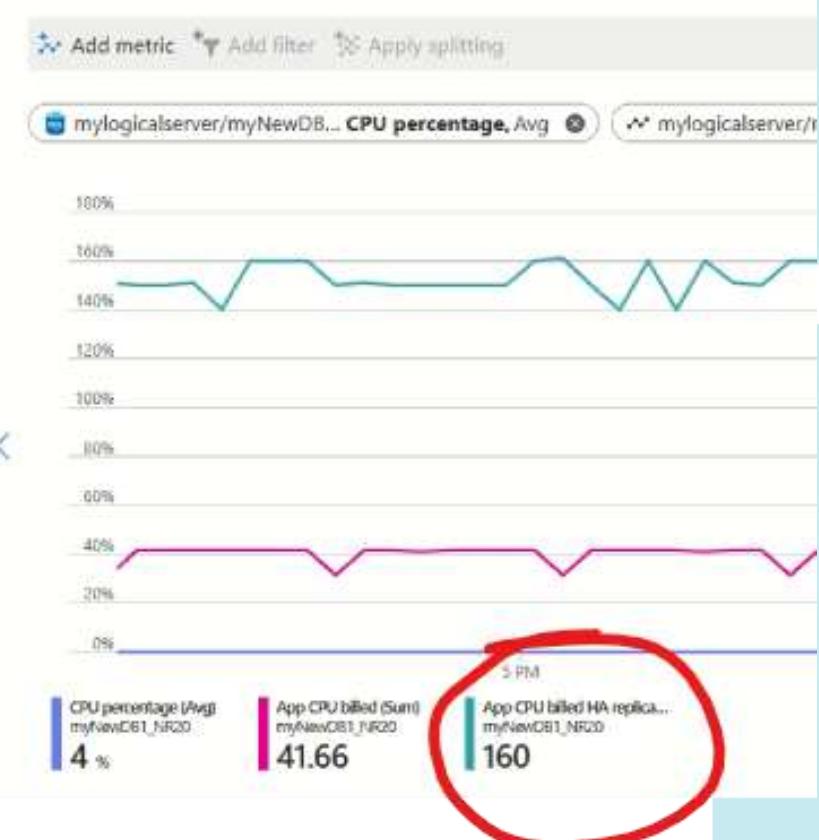
Replicas

Integrations

Azure Synapse Link Stream analytics (preview) Add Azure AI Search

Quiz time!

Why is the *App CPU billed (HA replicas)* so much higher than *App CPU billed*?



myNewDB1_NR20 (mylogicalserver/myNewDB1_NR20) | Metrics

X

SQL database



Search



+ New chart

⟳ Refresh

↗ Share

😊 Feedback

Local Time: Last hour (Automatic - 1 minute)

Data Discovery & Classification

Dynamic Data Masking

Microsoft Defender for Cloud

Identity

Data Encryption

Intelligent Performance

Performance overview

Performance recommendations

Query Performance Insight

Automatic tuning

Monitoring

Alerts

Metrics

Avg CPU percentage and Sum App CPU billed for myNewDB1_NR20



Add metric

Add filter

Apply splitting

Line chart

Drill into Logs

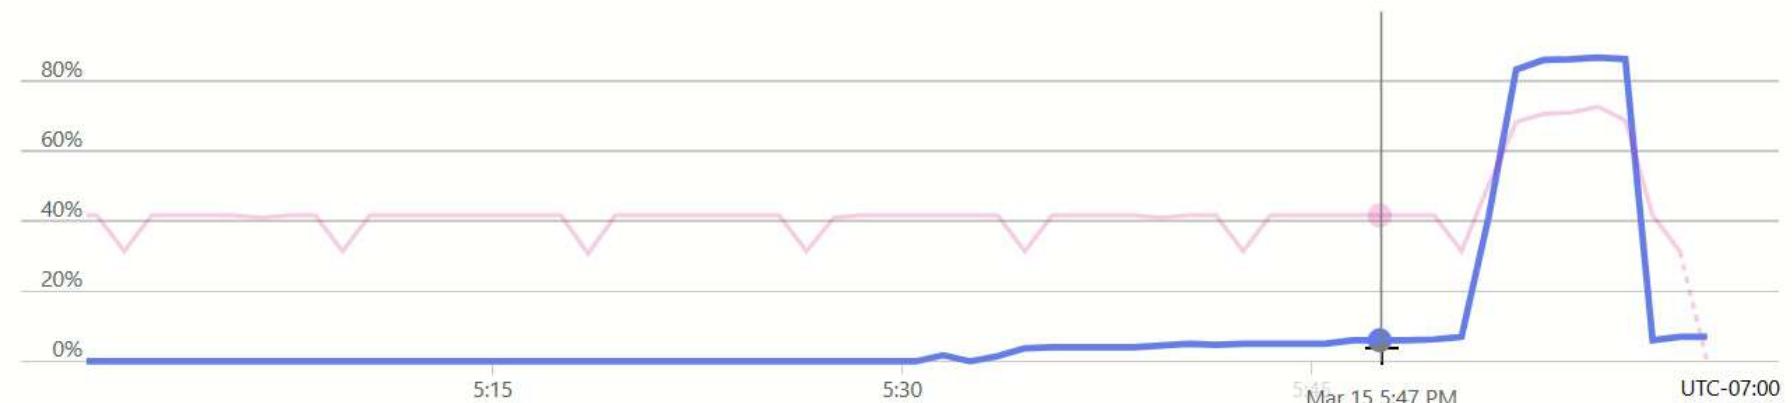
New alert rule

Save to dashboard

...

mylogicalserver/myNewDB... CPU percentage, Avg

mylogicalserver/myNewDB... App CPU billed, Sum



Writer @ 2024-03-16 01:19:10.963 Log rate 9.13 MB/sec; writer refresh 300 ms (press +/- to adjust)

Rdr # 0 OK @ 2024-03-16 01:19:09.635 Avg. latency: 59 ms	Rdr # 5 OK @ 2024-03-16 01:19:09.291 Avg. latency: 61 ms	Rdr # 10 OK @ 2024-03-16 01:19:09.635 Avg. latency: 54 ms	Rdr # 15 OK @ 2024-03-16 01:19:09.635 Avg. latency: 60 ms	Rdr # 20 OK @ 2024-03-16 01:19:09.291 Avg. latency: 55 ms	Rdr # 25 OK @ 2024-03-16 01:19:09.635 Avg. latency: 52 ms
Rdr # 1 OK @ 2024-03-16 01:19:09.635 Avg. latency: 56 ms	Rdr # 6 OK @ 2024-03-16 01:19:09.291 Avg. latency: 59 ms	Rdr # 11 OK @ 2024-03-16 01:19:09.635 Avg. latency: 52 ms	Rdr # 16 OK @ 2024-03-16 01:19:09.635 Avg. latency: 50 ms	Rdr # 21 OK @ 2024-03-16 01:19:09.635 Avg. latency: 51 ms	Rdr # 26 OK @ 2024-03-16 01:19:09.635 Avg. latency: 55 ms
Rdr # 2 OK @ 2024-03-16 01:19:09.635 Avg. latency: 57 ms	Rdr # 7 OK @ 2024-03-16 01:19:09.291 Avg. latency: 56 ms	Rdr # 12 OK @ 2024-03-16 01:19:09.635 Avg. latency: 53 ms	Rdr # 17 OK @ 2024-03-16 01:19:09.291 Avg. latency: 61 ms	Rdr # 22 OK @ 2024-03-16 01:19:09.635 Avg. latency: 50 ms	Rdr # 27 OK @ 2024-03-16 01:19:09.635 Avg. latency: 52 ms
Rdr # 3 OK @ 2024-03-16 01:19:09.635 Avg. latency: 51 ms	Rdr # 8 OK @ 2024-03-16 01:19:09.635 Avg. latency: 52 ms	Rdr # 13 OK @ 2024-03-16 01:19:09.635 Avg. latency: 51 ms	Rdr # 18 OK @ 2024-03-16 01:19:09.635 Avg. latency: 50 ms	Rdr # 23 OK @ 2024-03-16 01:19:09.291 Avg. latency: 57 ms	Rdr # 28 OK @ 2024-03-16 01:19:09.635 Avg. latency: 62 ms
Rdr # 4 OK @ 2024-03-16 01:19:09.635 Avg. latency: 51 ms	Rdr # 9 OK @ 2024-03-16 01:19:09.635 Avg. latency: 56 ms	Rdr # 14 OK @ 2024-03-16 01:19:09.635 Avg. latency: 53 ms	Rdr # 19 OK @ 2024-03-16 01:19:09.635 Avg. latency: 50 ms	Rdr # 24 OK @ 2024-03-16 01:19:09.635 Avg. latency: 60 ms	Rdr # 29 OK @ 2024-03-16 01:19:09.291 Avg. latency: 58 ms

Hyperscale Elastic Pools

Public preview



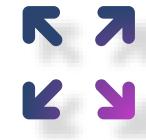
Increased storage limits

Storage up to 100 TB per elastic pool



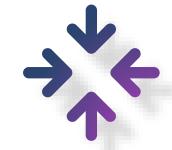
Higher log generation rates

Log throughput independent of compute



Predictable scaling

Scale the compute for the elastic pool up or down in minutes, regardless of how big the pooled DBs are!



Pool-level read scale

Optional pool replicas to host read-only workload

<https://aka.ms/hsepvideo> (4 minute video); <https://aka.ms/hsep> (docs)
<https://aka.ms/hsep-public-preview-tech-blog> (public preview announcement).

Energy



Payment Solutions



Health Care



Entertainment



Others



Deutsche Post DHL
Group

Takeaways

Hyperscale is THE cloud native database for all types of workloads.

Hyperscale grows as your workloads grow – offering great scalability.

Hyperscale offers leading price-performance, especially with options like serverless and elastic pools.

Please do provide feedback!

<https://sqlb.it/?12734>

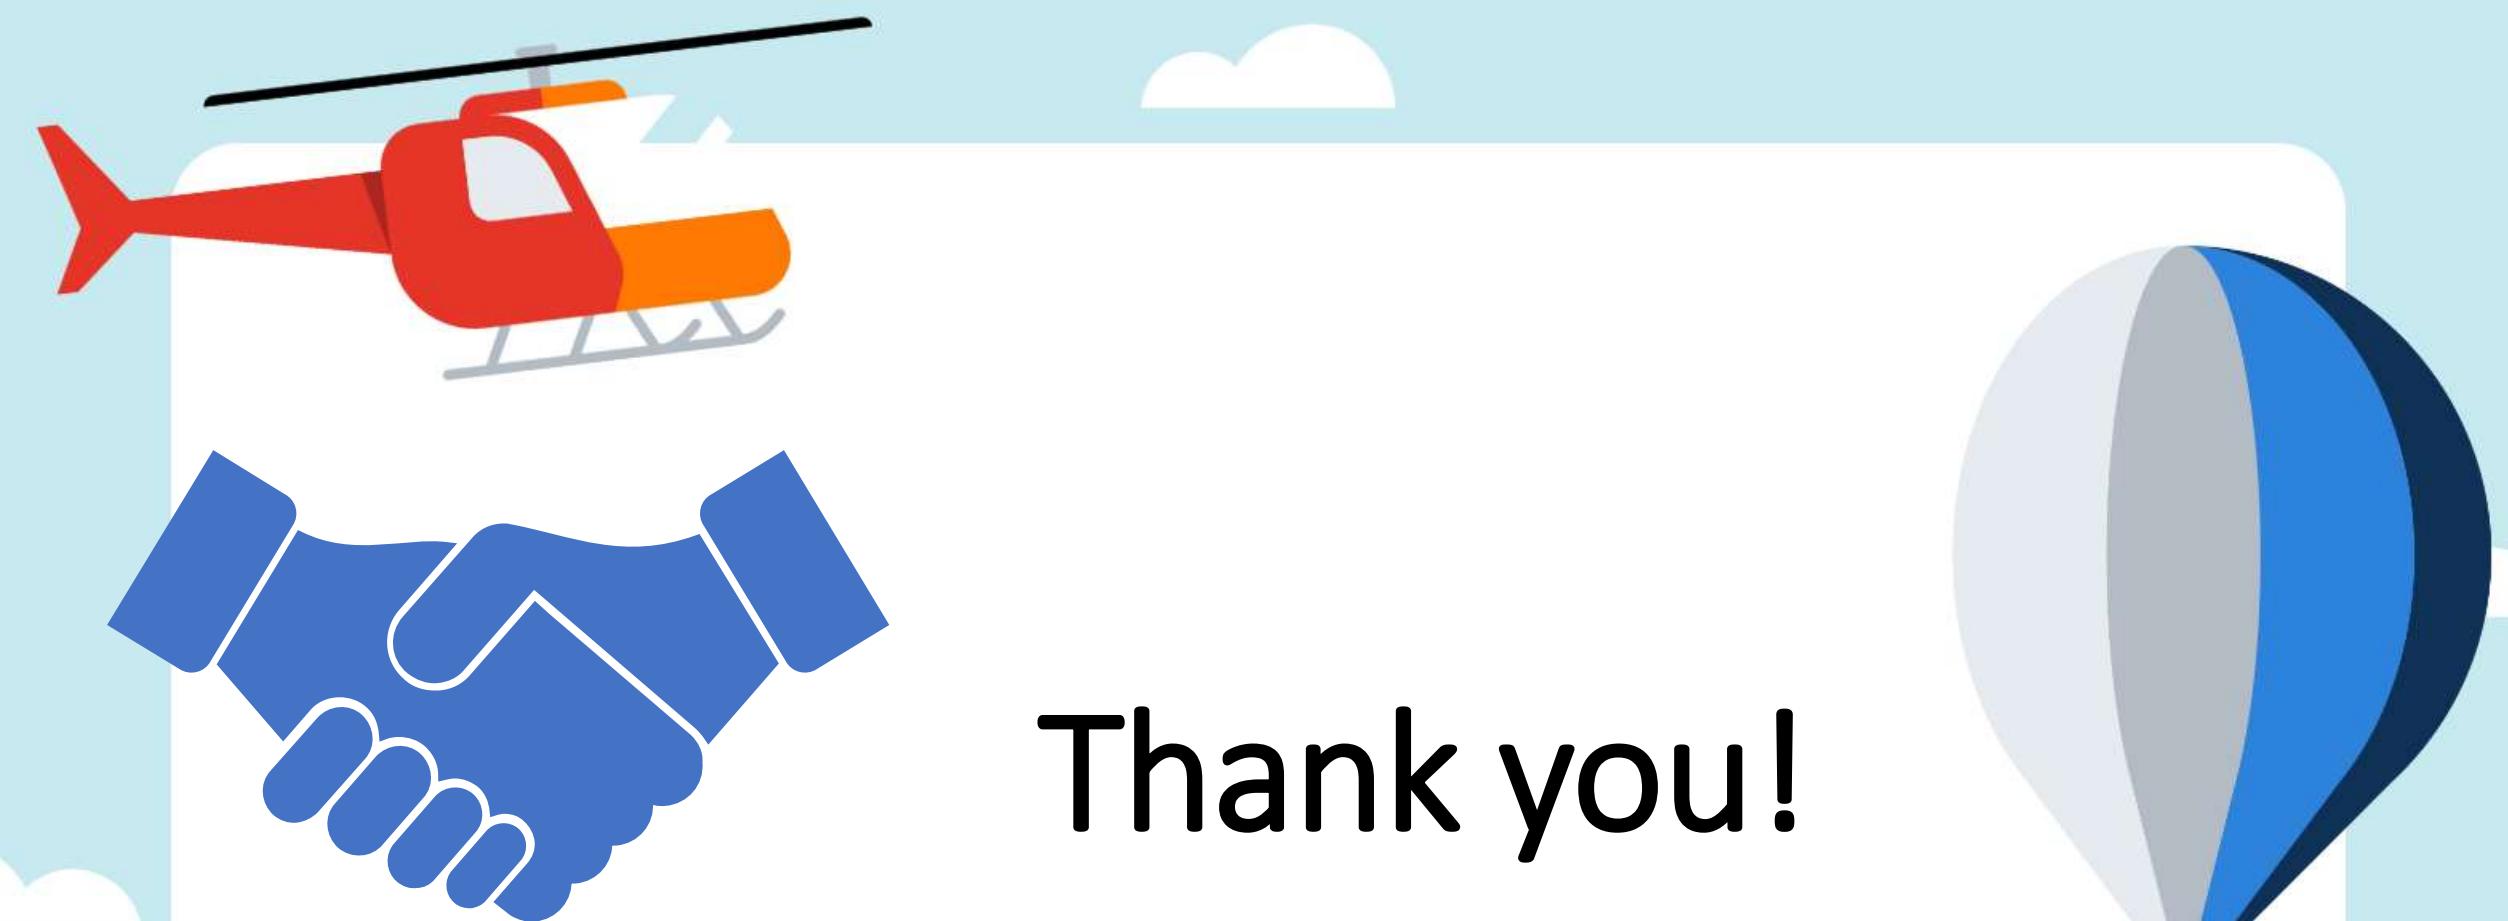


References



- More about Hyperscale: <https://aka.ms/hs>
- Hyperscale simplified and lower pricing:
<https://aka.ms/hsignite2023> and
<https://aka.ms/hsprice2023>
- Hyperscale serverless:
<https://aka.ms/sqlserverless>
- Hyperscale elastic pools:
<https://aka.ms/hsep-prms> and
<https://aka.ms/hsep>





Thank you!

@aditya_feb22 (X)

@arvisam (Bluesky, Mastodon and X)

