




Flying High Amongst the Clouds with Fabric Warehouses

Andy Cutler & Brian Bønck



Agenda – part 1

- Welcome and who are we
- What is Fabric an introduction
- Fabric Data warehouse vs other data services
- Architectures and architectural patterns
- The Polaris engine
- OneLake
- The T-SQL surface area
- Isolation levels
- Ingestion – load patterns from external sources



Agenda – part 2

- CRUD
- Table cloning
- Meta-data
- Security
- Disaster recovery and business continuity
- Source control and deployment
- Semantic Models – both default and user defined
- Direct Lake for DWH vs import
- Round up and Q/A



Who are we

Andy Cutler

Brian Bønck





sqlbits
2024



sqlbits
2024



Get data



Data Factory



Synapse Data
Engineering



Synapse Data
Warehouse



Synapse Data
Science



Synapse Real-
Time Analytics

Use data



Power BI




Data Activator
(coming soon)

Store data



OneLake





VS



VS



VS




Synapse Data Warehouse

Synapse Serverless SQL

Synapse Dedicated pool

Databricks





VS



Synapse
Serverless SQL

VS



Synapse
Dedicated pool

VS




Databricks

Storage: Fabric focuses on a single storage solution, the Unified Data Location, built on Azure Data Lake Storage Gen2. This eliminates the need for role-based access control and simplifies data management for organizations.

Integration: Seamless integration with multiple Azure services, including Azure Data Factory, Azure Synapse Analytics, Synapse Real-Time Analytics, Synapse Data Warehousing, Power BI and Data Explorer, with automatic provisioning of the underlying hardware.

Collaboration: Dedicated workspaces enable diverse developers – including data engineers and data scientists – to collaborate effortlessly.





Synapse Data
Warehouse

VS



Synapse
Serverless SQL

VS



Synapse
Dedicated pool

VS




Databricks

A unified workspace: Synapse Analytics provides a unified workspace where data engineers and data scientists can collaborate on big data and SQL-based analytics tasks.

Real-time analysis: The integration of Apache Spark and dedicated SQL pool enables users to perform real-time analysis - on both structured and unstructured data.

Security and Governance: Synapse Analytics provides robust security features and fine-grained access controls to ensure data is protected.





Synapse Data
Warehouse

VS



Synapse
Serverless SQL

VS



Synapse
Dedicated pool

VS



Databricks

Scalability: Databricks can easily handle large data tasks as it has the ability to scale horizontally.

Notebooks: Interactive notebooks allow users to execute and visualize code in real time, facilitating data exploration and model training.

Integration: Seamless integration with other Azure services such as Azure Storage and Azure Data Lake Storage simplifies the process of data import and extraction.



Data Warehouse

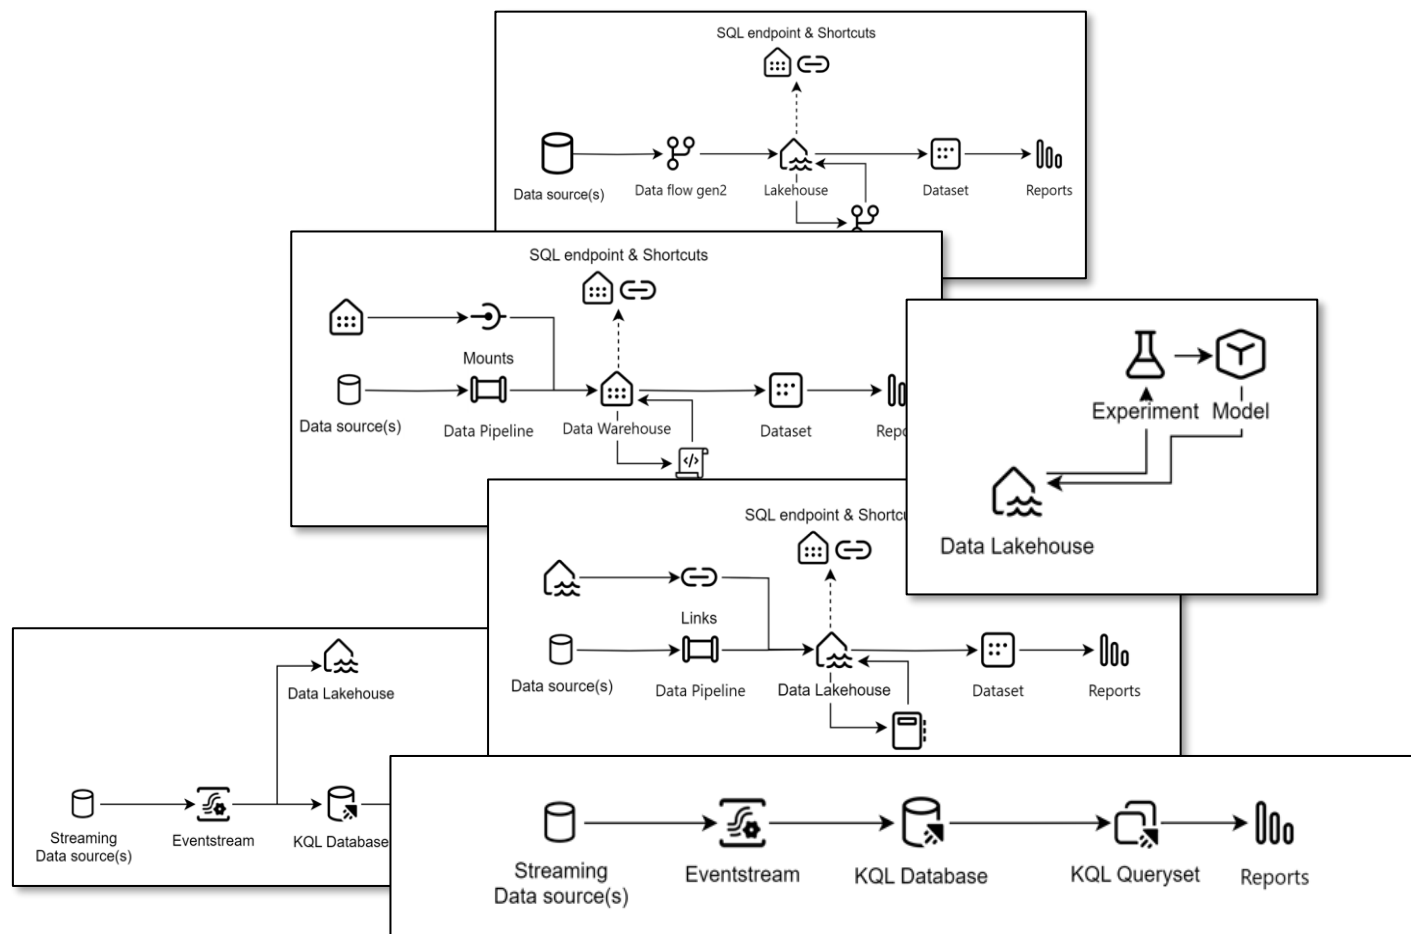


Lakehouse

Warehouse	Lakehouse
Structured	Structured, Semi structured, Unstructured
SQL language	Notebooks
Schemas, Tables	Folders, Files and tables
OLS, RLS, CLS	RLS, TLS
T-SQL	Spark (Scala, PySpark, SparkSQL)

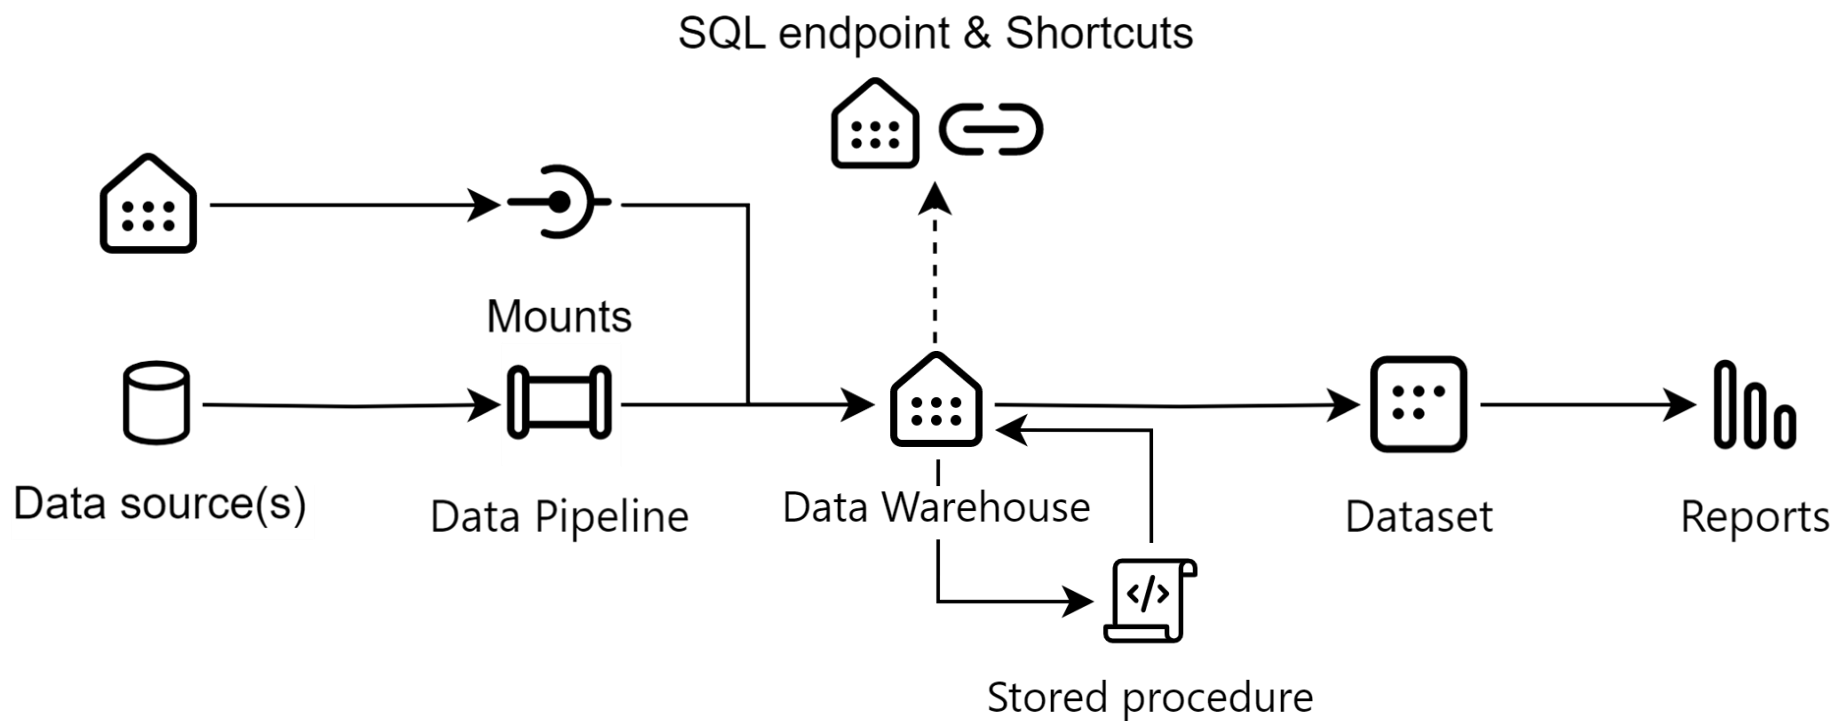


Data
Warehouse





Data
Warehouse





Data
Warehouse

Data Ingestion




Data Engineering




Power BI




Data Source



Raw



Cleaned



Transformed




Dataset




Report



Data
Warehouse

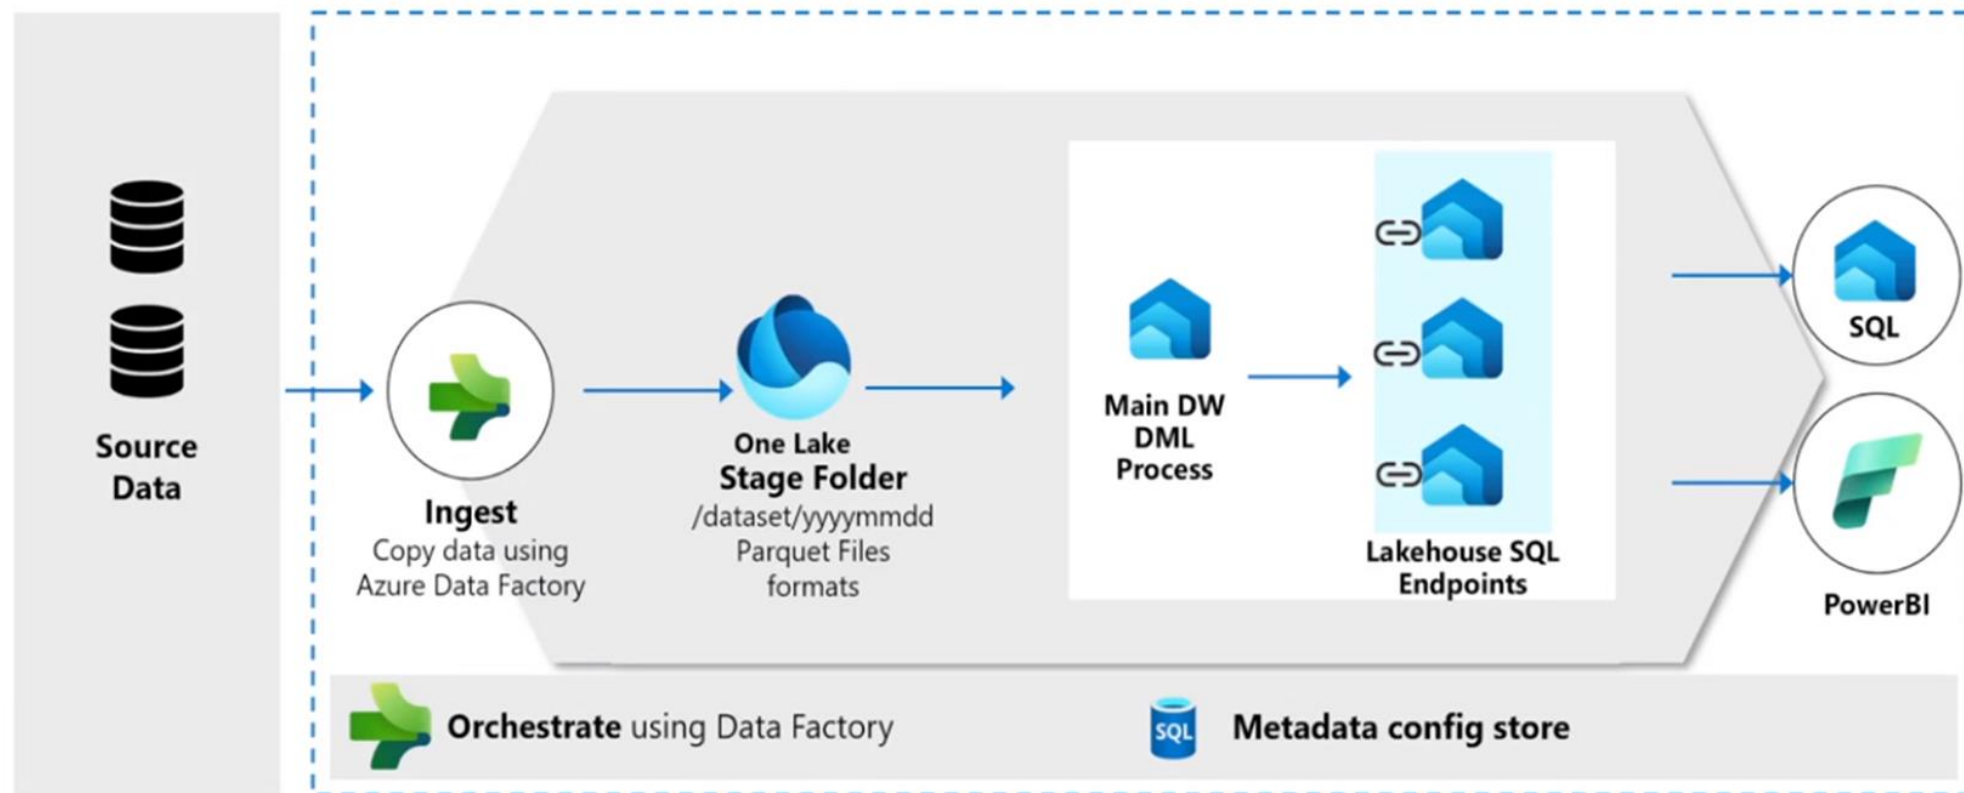


Logical Warehouse *Viewhouse*






Modern Data Warehouse Pattern in Microsoft Fabric



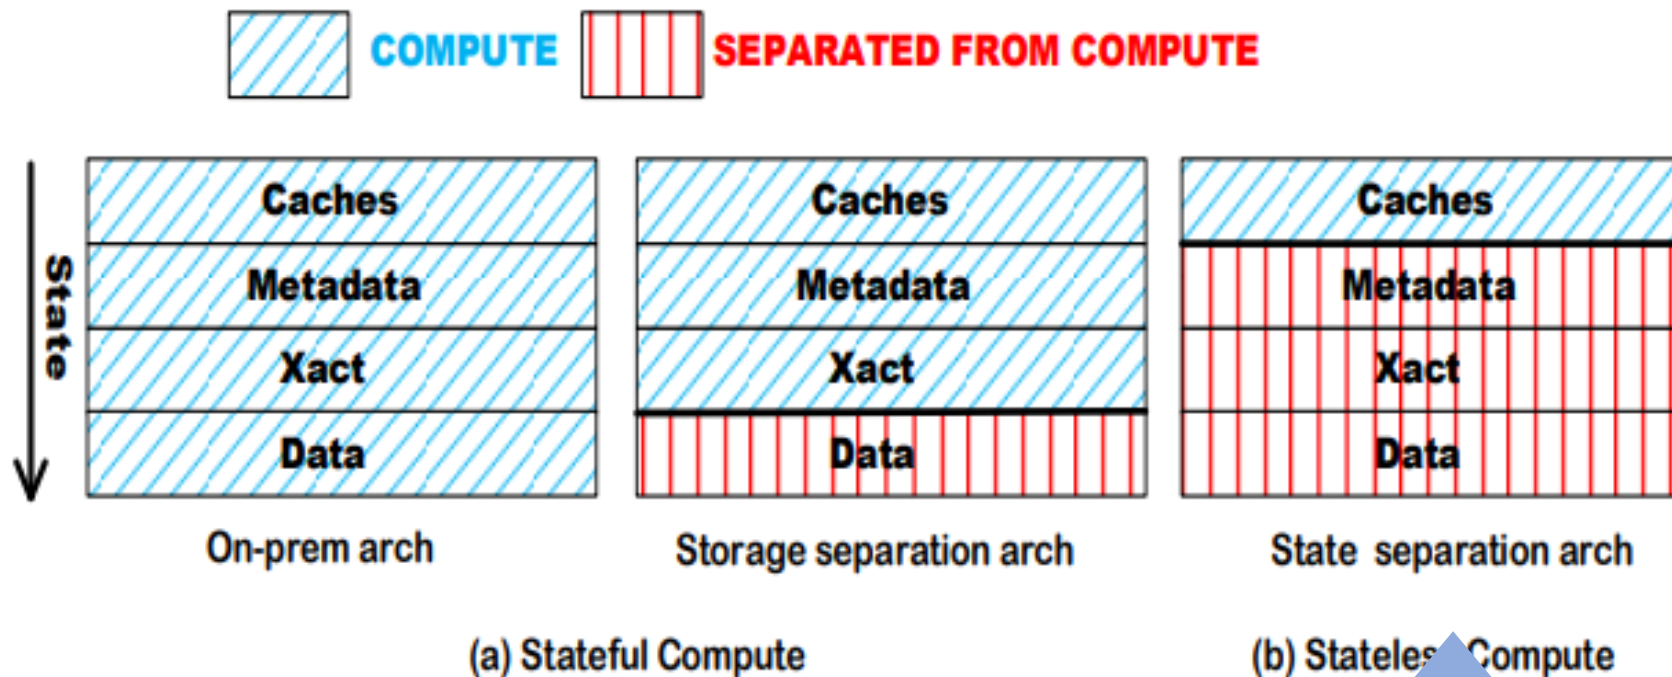


Data
Warehouse



The Polaris engine

Stateless compute



Synapse and Fabric






The Polaris engine

Data abstraction



Data
Warehouse

Data






The Polaris engine

Data abstraction



Data
Warehouse

User
partitions






The Polaris engine

Data abstraction



Data
Warehouse



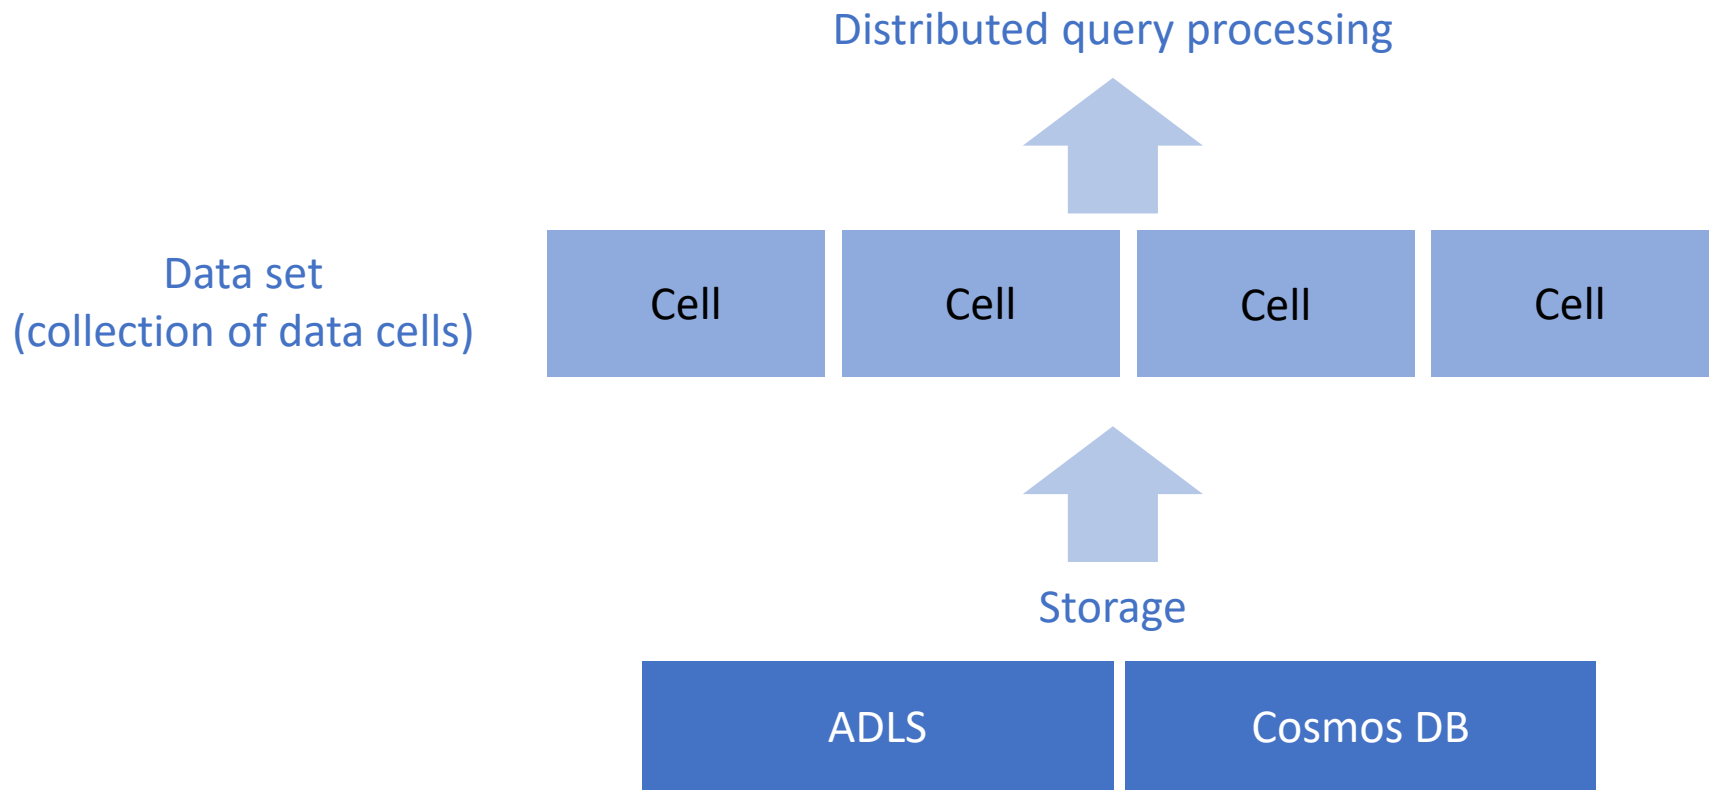


The Polaris engine

Data abstraction



Data
Warehouse






The Polaris engine

Data abstraction



Data
Warehouse



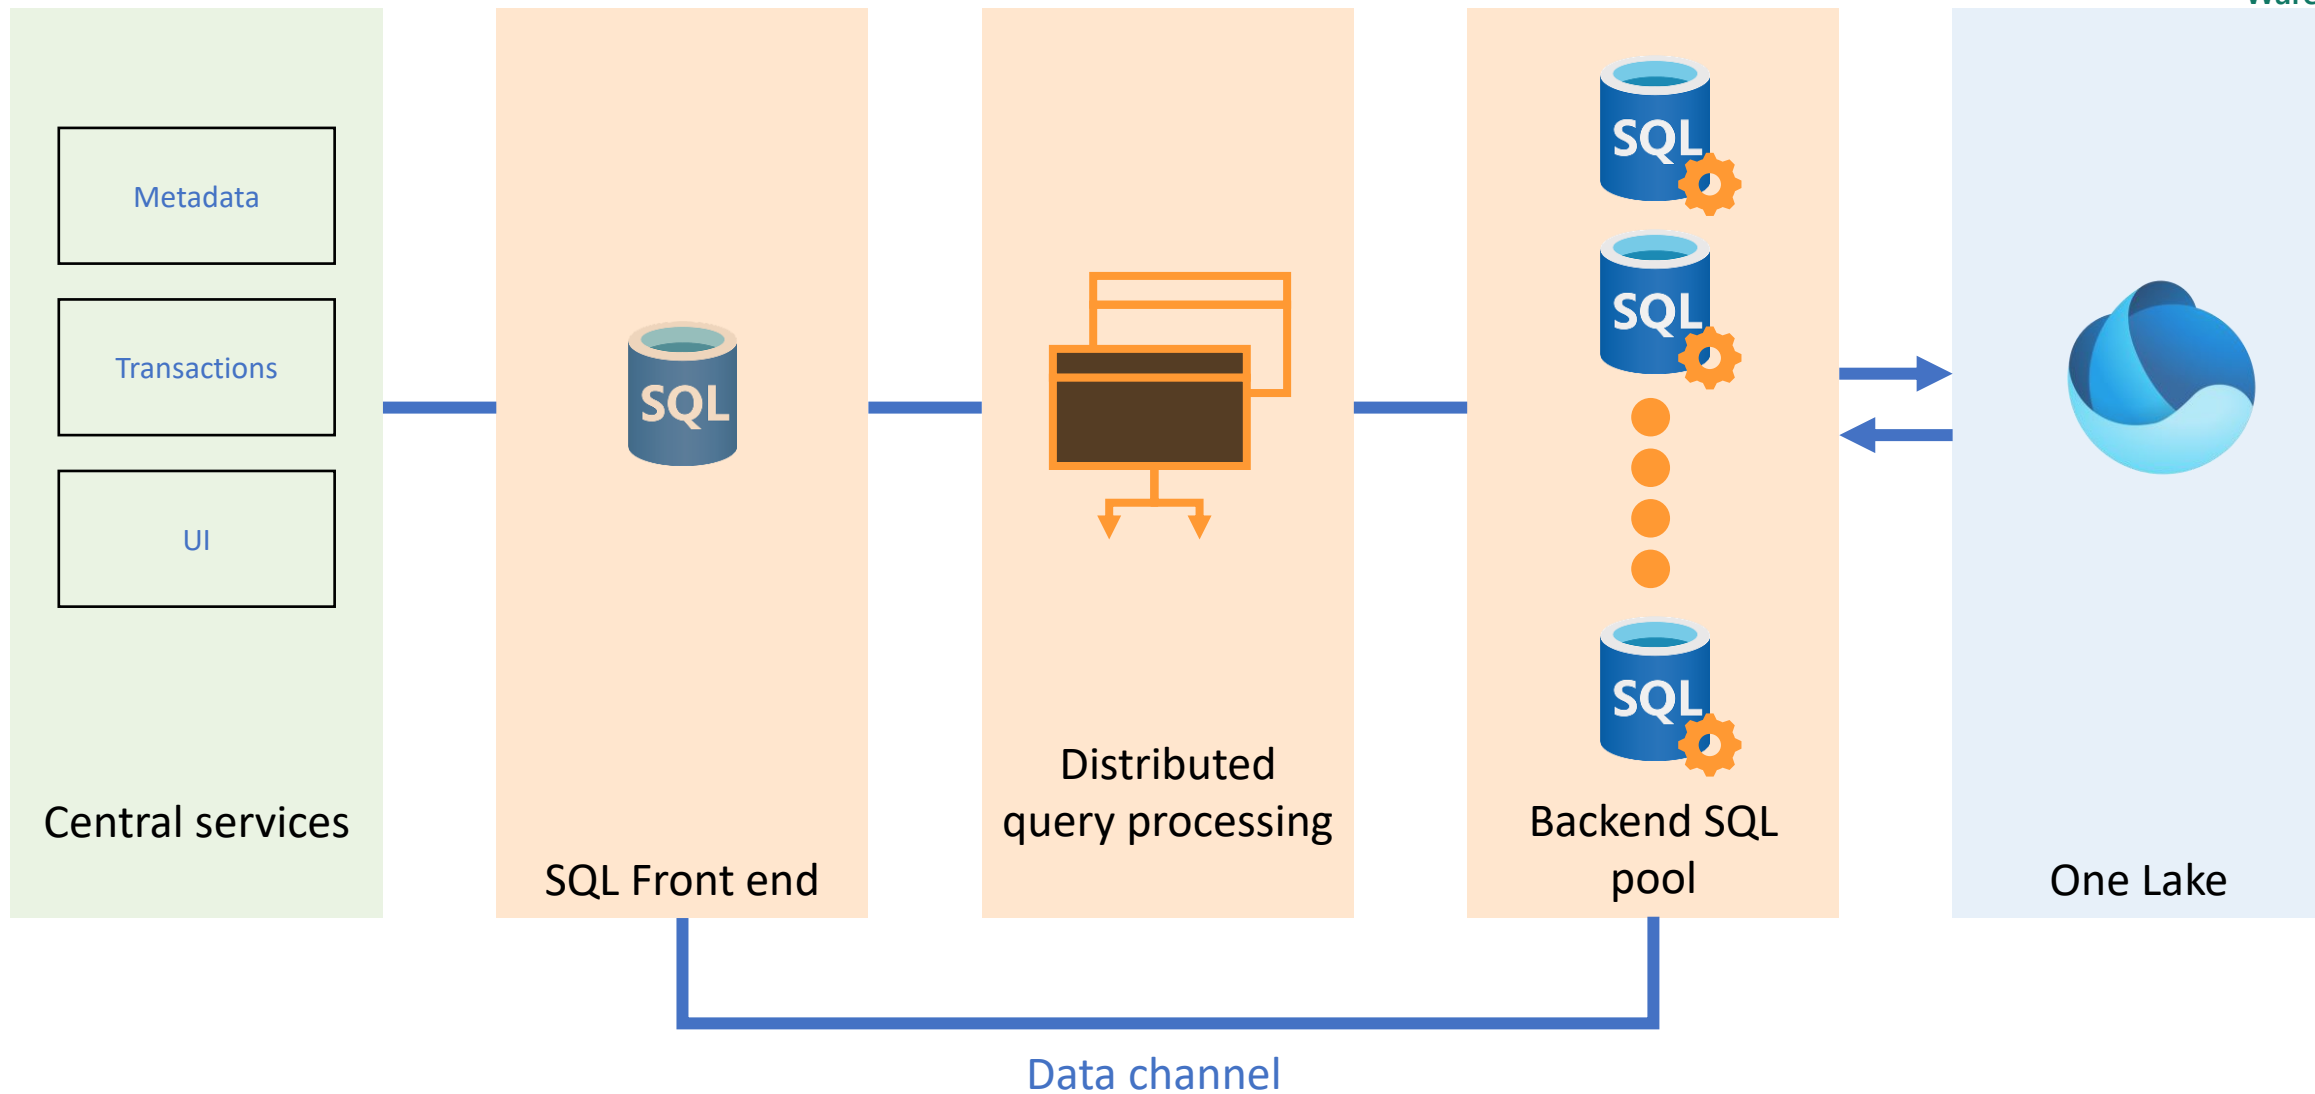


The Polaris engine

Service architecture



Data
Warehouse





The Polaris engine

Read more here



Data
Warehouse



Whitepaper on Polaris:
aka.fm/polaris



Broken down to a blogpost:
aka.fm/polaris-blog



T-SQL Surface Area



The T-SQL Surface area covers the supported syntax within both the Warehouse & the Lakehouse SQL Endpoint

Currently unsupported syntax (and the potential impact...)



ALTER TABLE...
...<ADD/ALTER/DROP>

Identity Columns

TRUNCATE



CREATE ROLE/USER

MERGE

Recursive Queries



BULK LOAD

SP_SHOWSPACEUSED

SET TRANSACTION
ISOLATION LEVEL

Transactions & Isolation Levels



Transactions are included (including multi-table)

Only **Snapshot Isolation** supported (Optimistic Concurrency)

Locks are at the **table level**

Use **sys.dm_tran_locks** to show current locks

Schema Stability
Sch-S

SELECT

Intent Exclusive
IX


INSERT
DELETE
UPDATE
COPY INTO

Schema
Modification
Sch-M

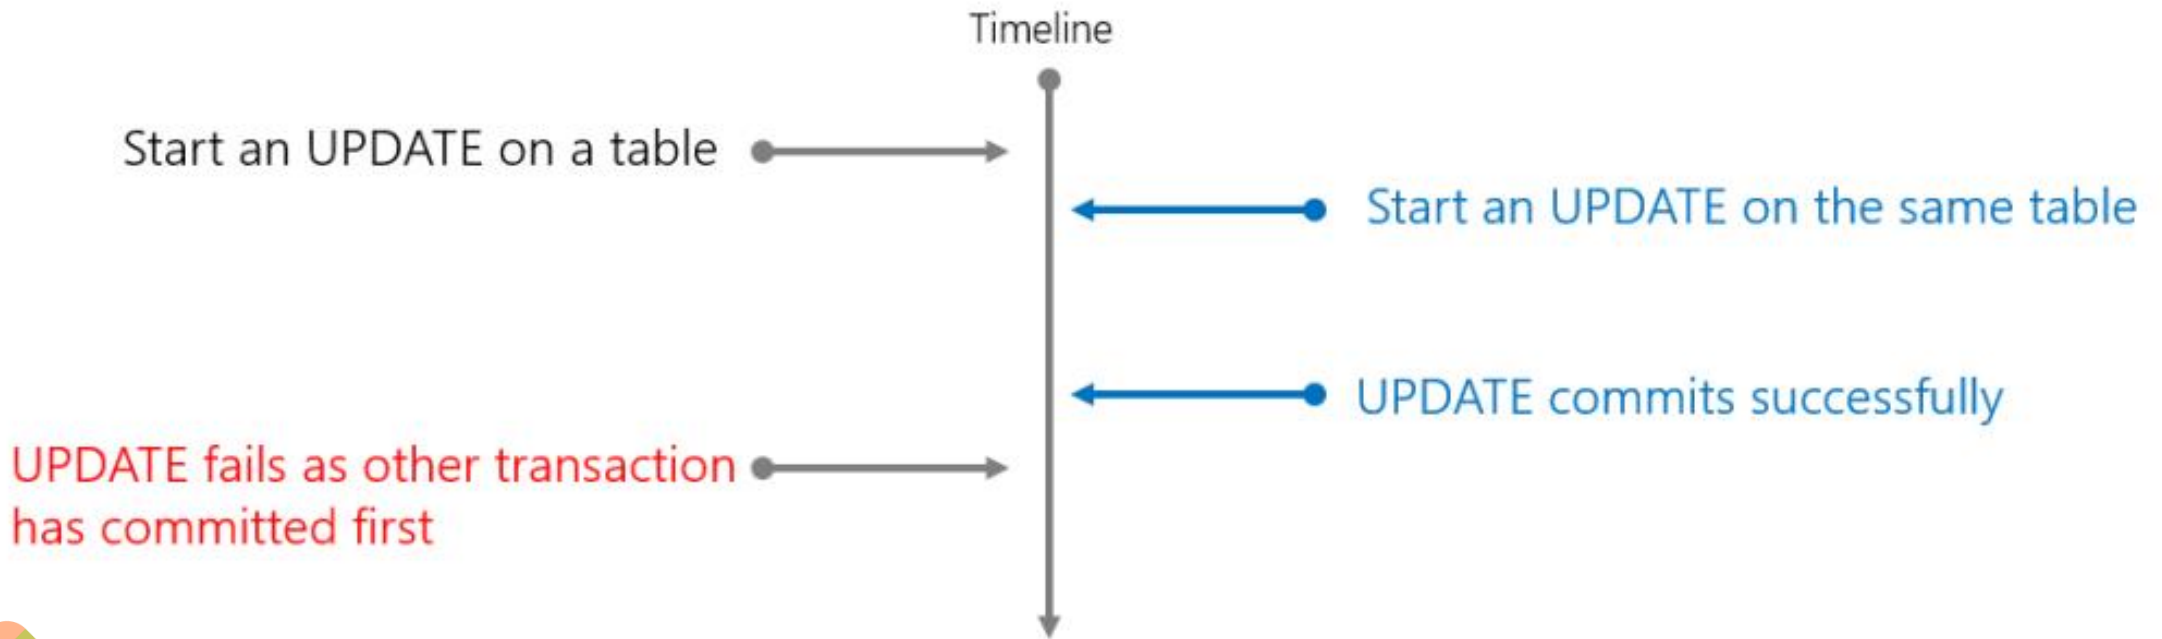
DDL



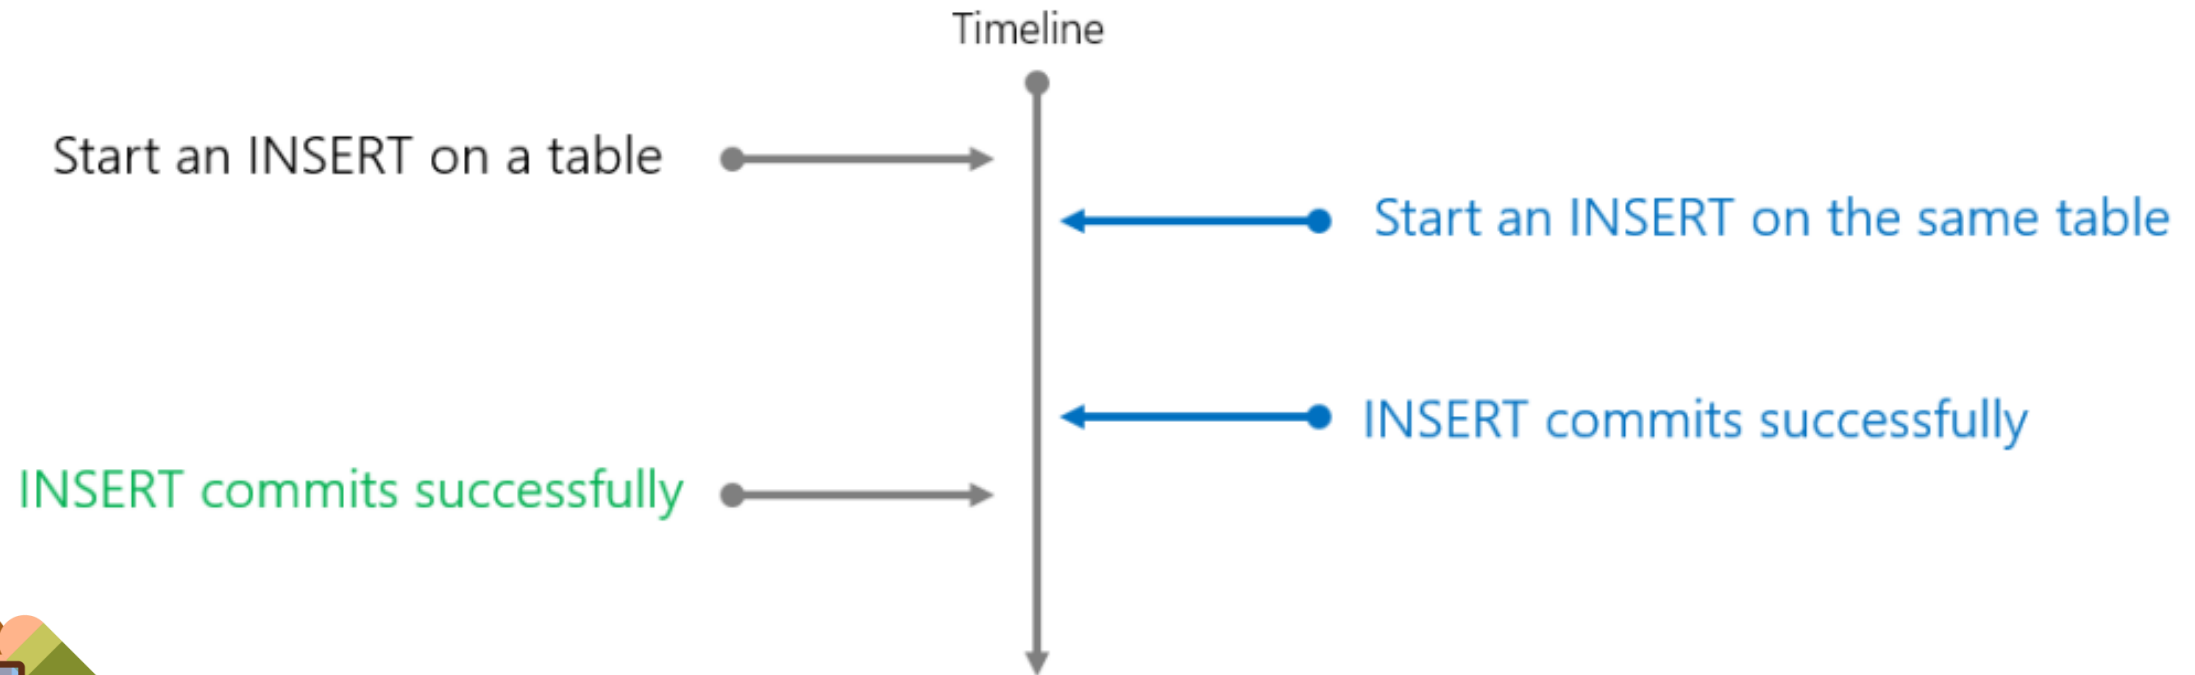
Transactions & Isolation Levels



Transactions & Isolation Levels



Transactions & Isolation Levels





Ingestion: Loading the Warehouse



Data
Warehouse

Code

SQL **COPY INTO**... <from Azure storage>

SQL **CREATE TABLE AS...SELECT**

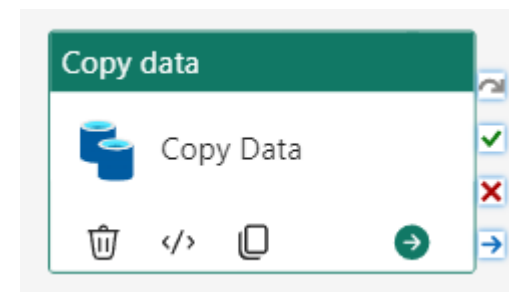
SQL **INSERT INTO...VALUES / SELECT**

Pyodbc/JDBC (from Notebooks)

Low/No-Code





Pipelines

“Brute force”



Dataflows Gen2
(Power Query)

Update method



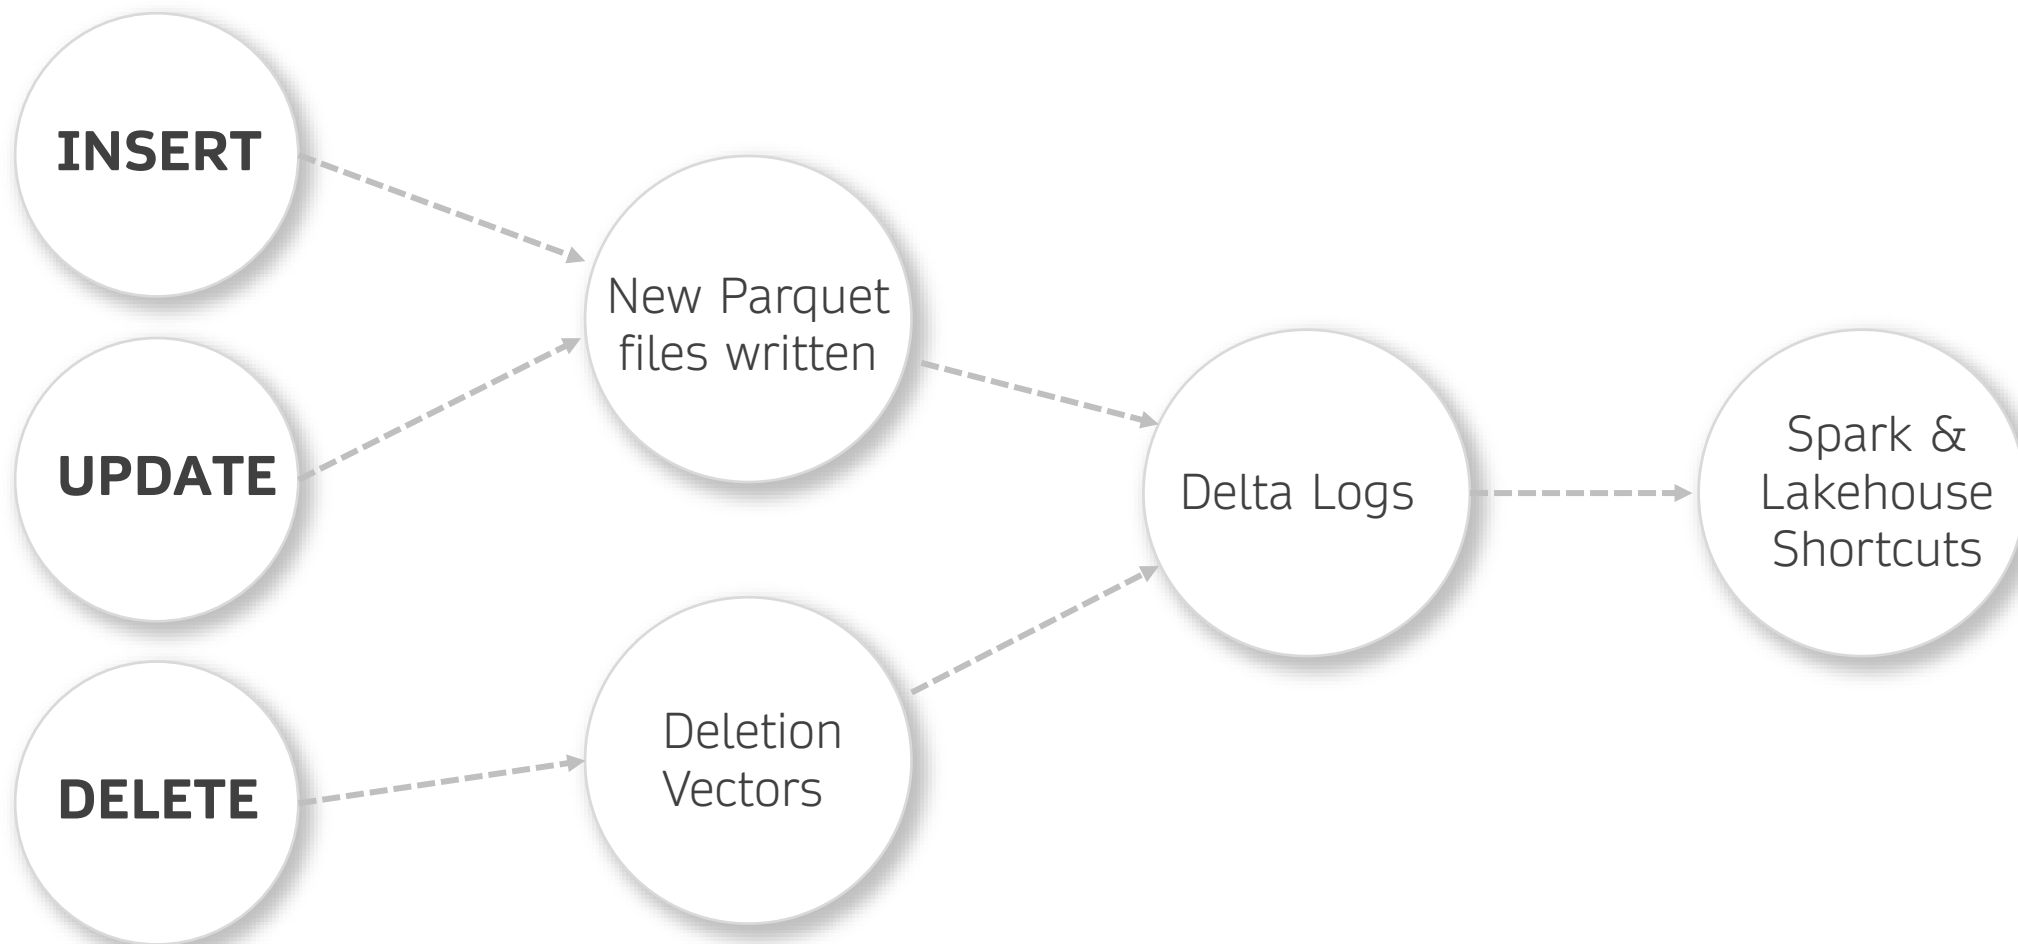


Ingestion: Loading the Warehouse



Data
Warehouse

All SQL operations are performed on Parquet files in OneLake





Ingestion: Loading the Warehouse

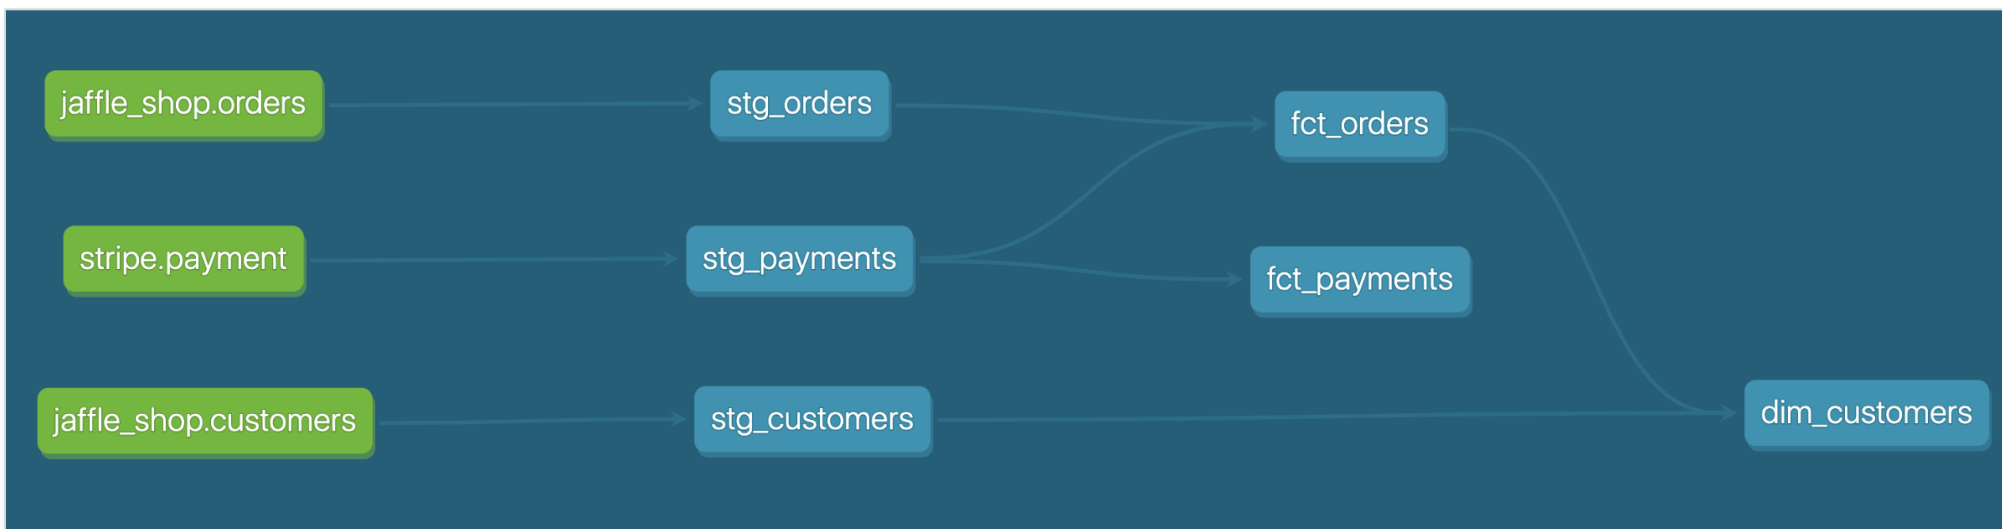


Data
Warehouse

dbt (including cloud) support in Warehouse

Declare your SQL loading processes

Create DAG for loading tables based on dependencies



Warehouse Utilization




Data
Warehouse

Usage is tracked by Capacity Unit seconds (consumed by read and write activity against the Warehouse and reads against Lakehouse SQL Endpoint)

- Warehouse Query: Compute charged for Warehouse (includes user generated and system tasks)
- SQL Endpoint Query: Compute charged for Lakehouse SQL Endpoint (includes user generated and system tasks)
- One Lake compute: Compute charged for all reads and writes for data stored in One Lake

Operation name	CU (s)	Duration (s)	Users	Billing type
Warehouse Query	2,187.77	1,360.42	5	Billable
OneLake Compute	0.01	11,880.00	1	Billable
Total	2,187.78	13,240.42	6	

Operation name	CU (s)	Duration (s)	Users	Billing type
SQL Endpoint Query	4,086.87	1,503.81	6	Both
Total	4,086.87	1,503.81	6	





Bursting and Smoothing

- Regardless of SKU, Fabric *bursting* will automatically allocate resources as needed to execute at maximum performance
- As such, one query could consume all the quota of a single time window and much more!
- To avoid an overload, *smoothing* kicks in

Dynamics Management Views



Data
Warehouse

Monitor Connection, Session, and Request Status in SQL Analytics Endpoint and Warehouses

sys.dm_exec_connections

- Provides comprehensive information about active connections to the Fabric Warehouse SQL Engine
- Includes details such as session ID, client address, client port, protocol, and authentication method.

sys.dm_exec_sessions

- shows information about all active user connections
- Offers insights into resource utilization by each session, such as the number of active requests

sys.dm_exec_requests

- Provides detailed information about currently executing or waiting queries
- Includes details such as session ID, query text, start time, status

Member, Contributor, and Viewer can execute sys.dm_exec_sessions and sys.dm_exec_requests

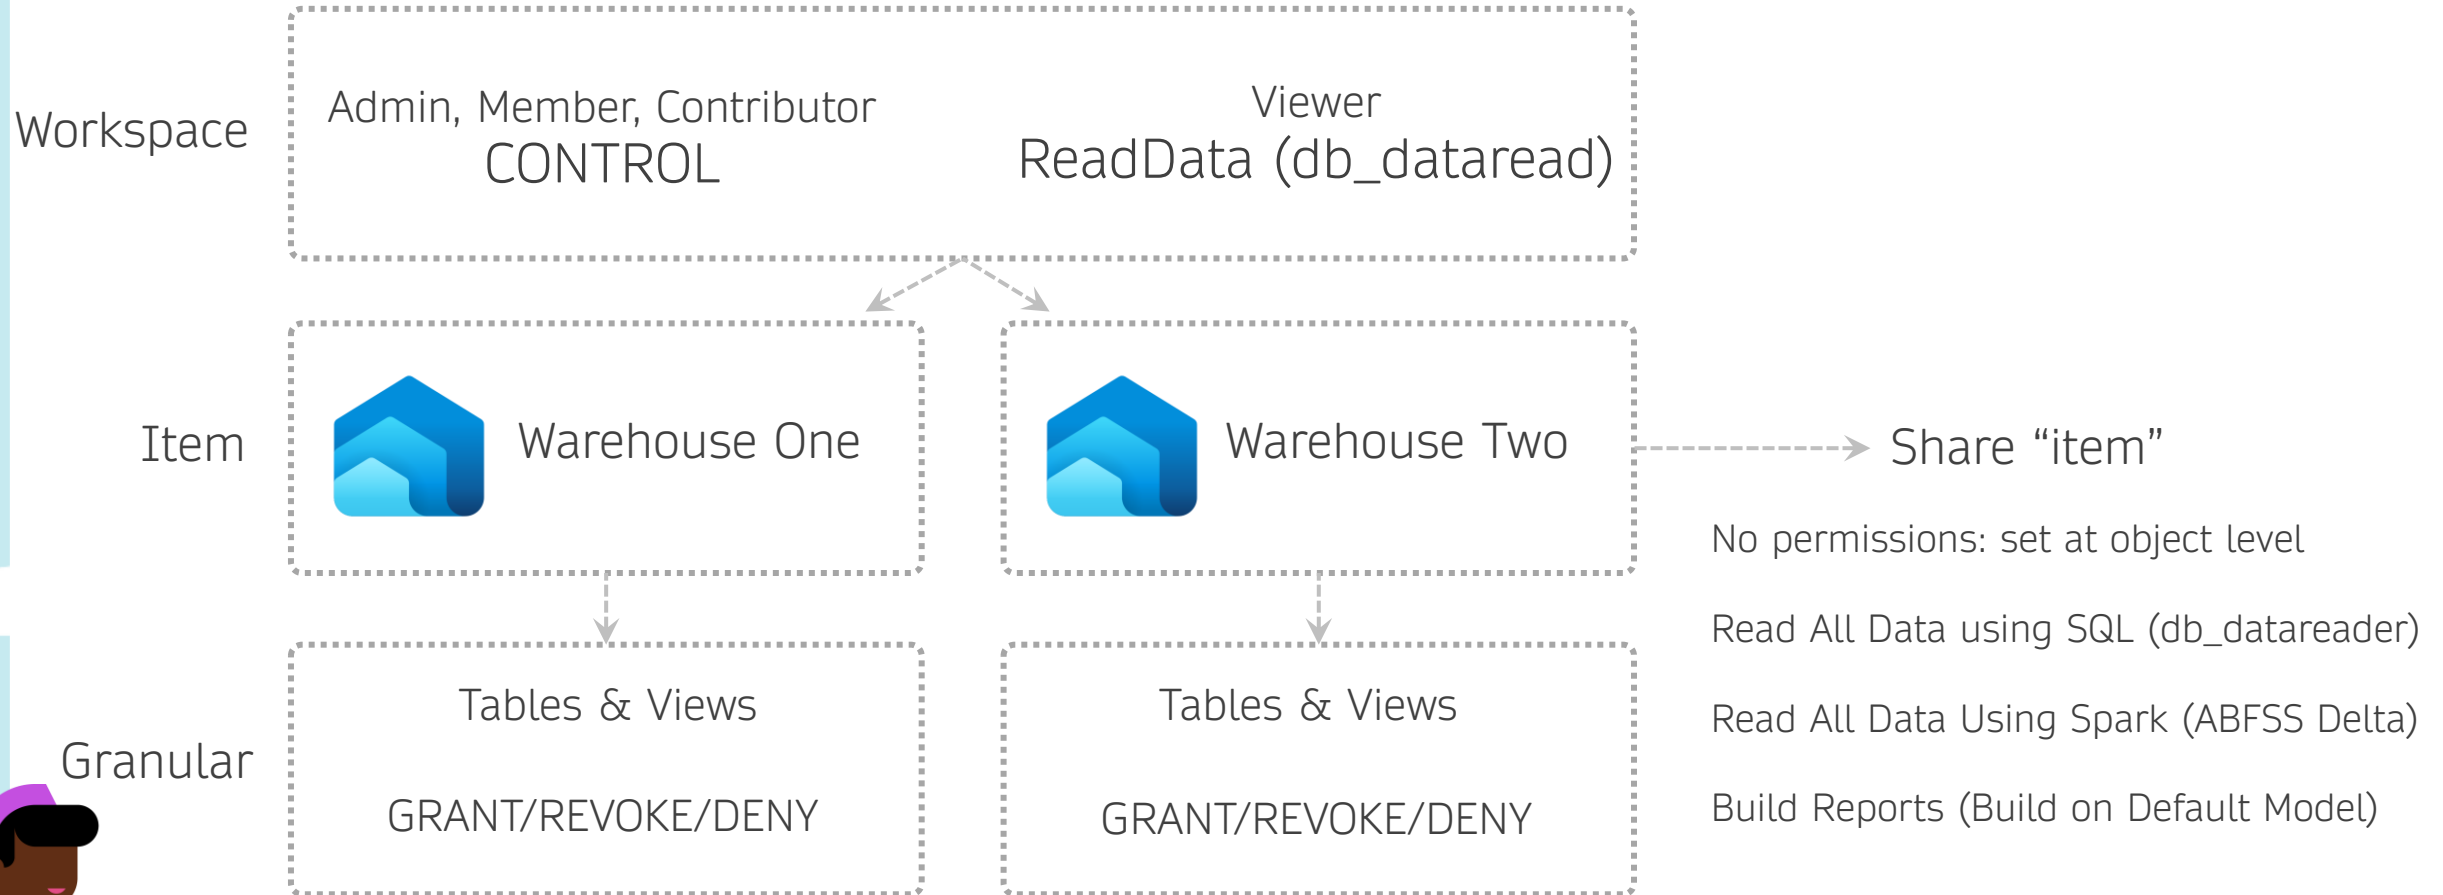


Security



Data
Warehouse

Security can be applied at Workspace, Warehouse, and Object level



Restore-In-Place

- Restore to a known “good” state
- Used when any corruption occurs
- Restore back after a failed deployment
- Restore back to a version for dev/test

Table Cloning (see next section)

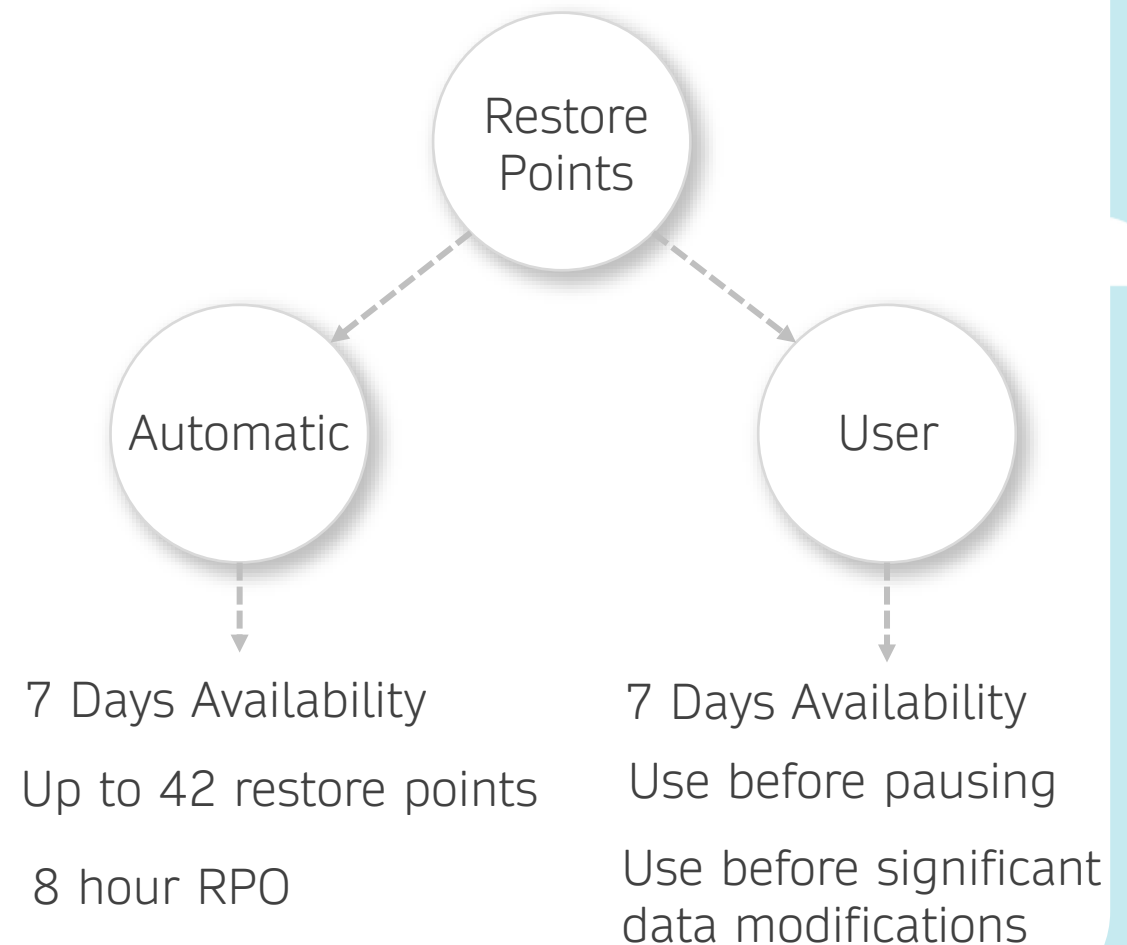





Table Cloning



Data
Warehouse

Clone an existing table into new table

“Shallow” clone as only the metadata is cloned

Cloned table is separate from base table

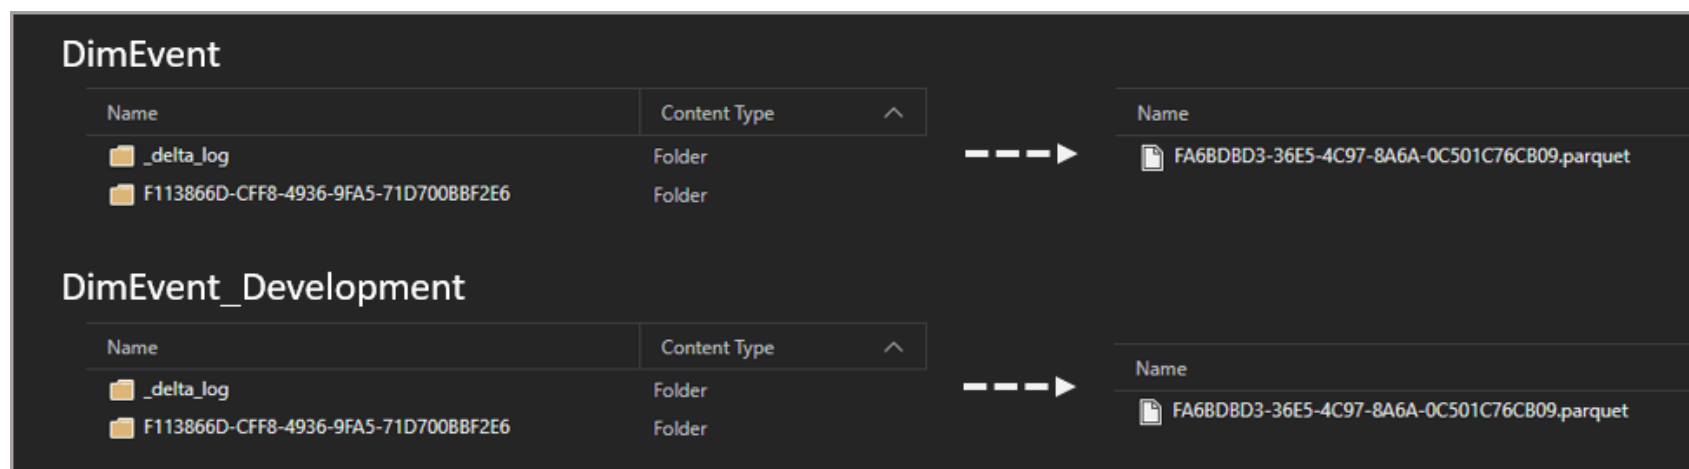
--create table clone

```
CREATE TABLE <schema>.<new_table_name>
```

```
AS CLONE OF <schema>.<existing_table_name>
```

Used for:

- Snapshotting at point-in-time
- Backups for recovery
- Testing data changes before applying to main table



Source Control & Deployment

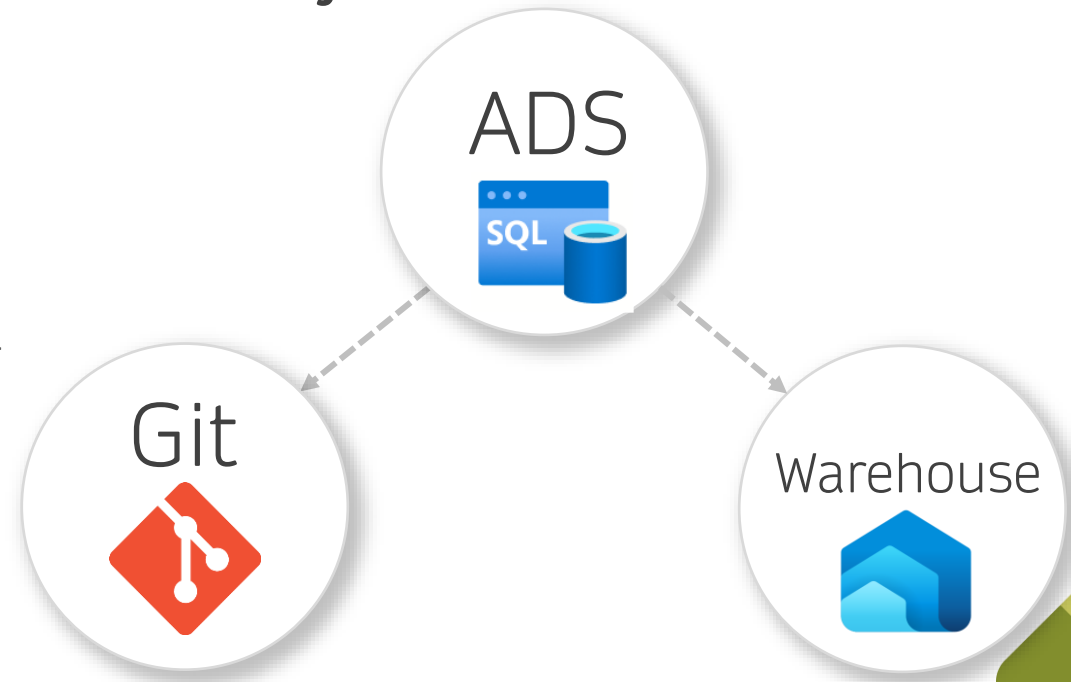


Data
Warehouse

“Treat your database like code” Grant Fritchey

Azure Data Studio Database Projects

- Current version 1.48.0 (DB: 1.4.2)
- Support for Warehouse and Lakehouse SQL Endpoint
- Uses dacpac for build and deploy
- Cannot ALTER!



Connect ADS to Git to
allow source control



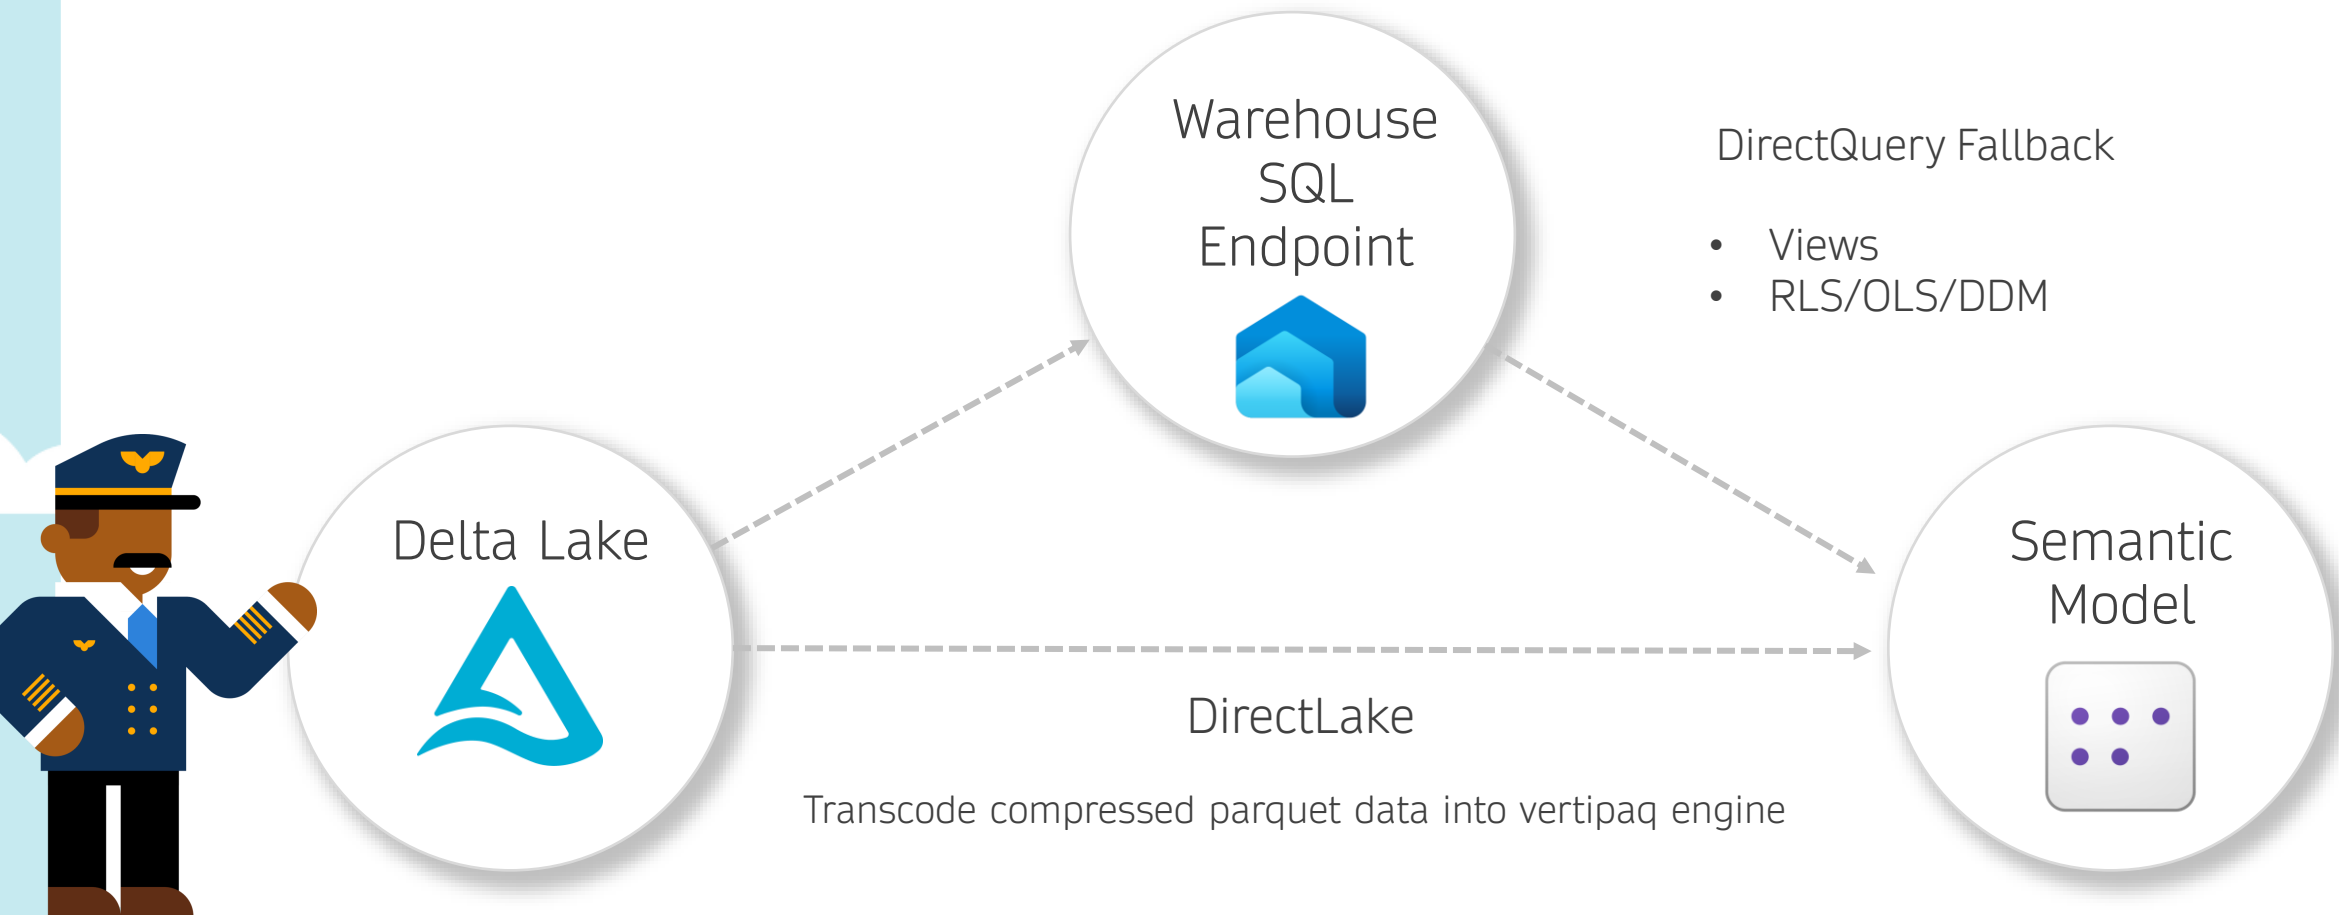
Semantic Models – Direct Lake



Data
Warehouse

Direct Lake is the new connectivity method from Semantic Models to Warehouse & Lakehouse

No requirement to Import the data into the Semantic model



Semantic Models – Relationships



Data
Warehouse

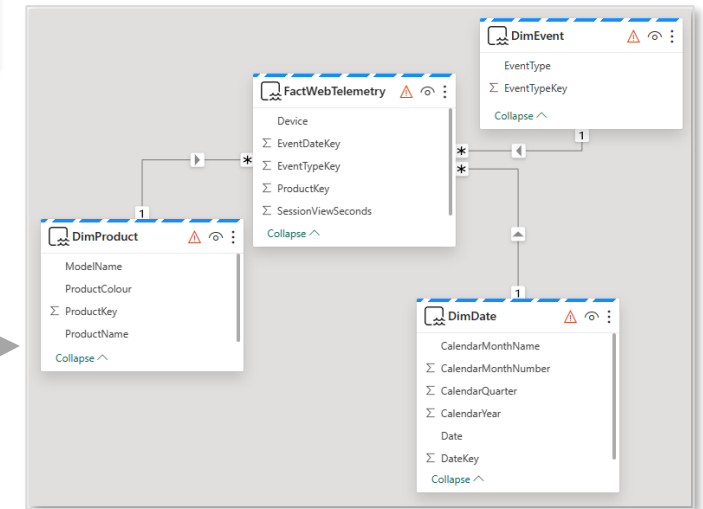
Constraints created in Warehouse appear in Default Semantic Model

Constraints created in Default Semantic Model appear in Warehouse



	ABC foreign_table	ABC primary_table	ABC fk_column	ABC pk_column ↑	ABC fk_constraint
1	dbo.FactWebTelemetry	dbo.DimDate	EventDateKey	DateKey	FK_FactWebTelemetry_EventDateKey
2	dbo.FactWebTelemetry	dbo.DimEvent	EventTypeKey	EventTypeKey	FK_FactWebTelemetry_EventTypeKey
3	dbo.FactWebTelemetry	dbo.DimProduct	ProductKey	ProductKey	FK_FactWebTelemetry_ProductKey

Default Semantic Model has a bi-directional sync with it's underlying Warehouse



Custom Semantic Models are unaffected by Warehouse constraints

Questions and dialogue

