

Fellowwind

From zero to Kusto hero at 30.000 ft

A unified analytics solution for the era of AI





Brian Bønk Rueløkke


Principal & Enterprise architect, Data & AI

Fellowmind

 <https://linkedin.com/in/brianbonk>

 <https://brianbonk.dk>

 <https://github.com/brianbonk>



AGENDA

The history of Kusto

Where does Kusto and RTA fit in the Data area

RTA in Fabric – incl. roadmap

Capabilities using Kusto

Get started for free

Introduction to the KQL language

BREAK

Kusto data in Power BI – with ninja tricks 

Kusto Functions

Next level KQL language

Outliers

Visualization

Dash-boarding



Jaques Cousteau 1910-1997



Jaques Cousteau 1910-1997

Image is AI generated

The history of Kusto



Azure Sentinel



Log Analytics



Real-Time Analytics




Azure resource
graph



Microsoft 365
Defender

CMPIvot

CMPIvot



The world is exploding with
high-granularity data...

It all starts with data

Telemetry – a key data for digital transformation

Telemetry – a key data for digital transformation




Telemetry – a key data for digital transformation

75B

Connected devices
by 2025

90ZB

Total amount of
telemetry data
by 2025



Digital transformation

Cybersecurity
Asset tracking and management
Predictive maintenance
Supply chain optimization
Customer experience
Energy management
Inventory management
Quality control
Environmental monitoring
Fleet management
Health and safety





Microsoft Fabric

Get data



Data Factory

Prepare data



Synapse Data
Engineering



Synapse Data
Warehouse



Synapse Data
Science



Synapse Real-Time
Analytics

Use data



Power BI



Data Activator

Store data



OneLake



Microsoft Fabric

Get data



Data Factory

Prepare data



Synapse Data
Engineering



Synapse Data
Warehouse



Synapse Data
Science



Synapse Real-Time
Analytics

Use data



Power BI



Data Activator

Store data



OneLake

Fabric Real-time Analytics solution enables organizations to consume **vast amount of data**, focus and **scale up** their Analytics solution with **data in motion**, **empower** their business analysts, and **democratize their data** for citizen data scientists and Data Engineers



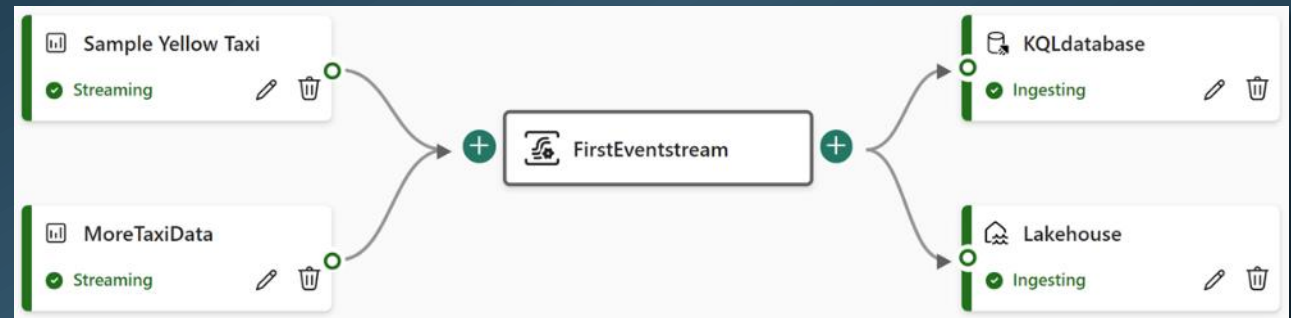
Synapse Real-Time
Analytics

⚡ Streaming data with ease



EVENTSTREAM

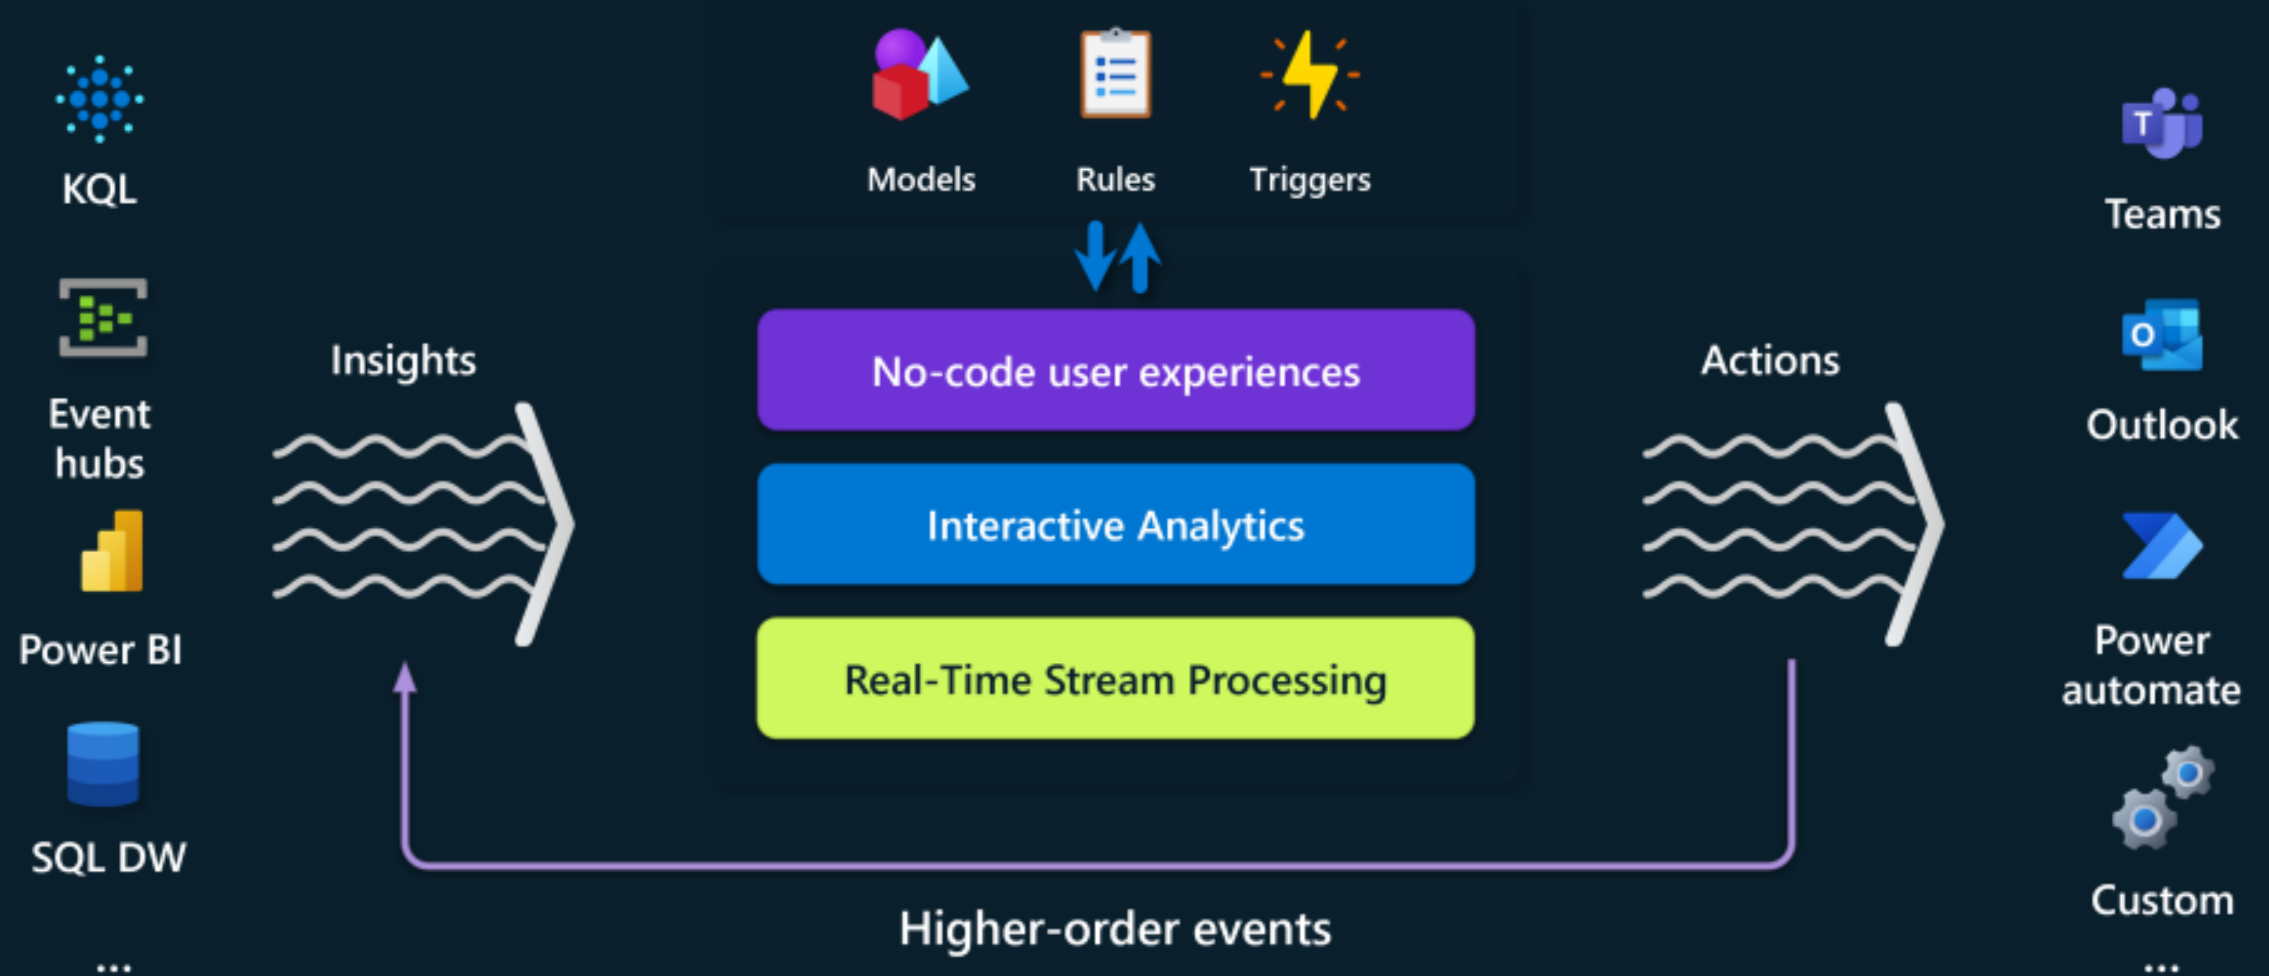
The brand-new event stream service, leverages the ability to get data from several sources of streaming data and save it to a wide variety of destinations, including OneLake, KQL databases and Azure services.



The service computes the data once and can pipe it out to several destinations at once. All configured and maintained from within the Microsoft Fabric portal and “coded” with your mouse.

Imagine scenarios of IoT devices loading data to both the data warehouse and other 3-rd party destinations – this can now be done using the low-code approach from Event Stream.

⚡ Data Activator



KQL database

Key capabilities

Unlimited Scale
(query, ingestion
and storage)

Any data source

Any data format

Structured
Semi-structured
Free-text

Real-time
transformation of
complicated data
structures

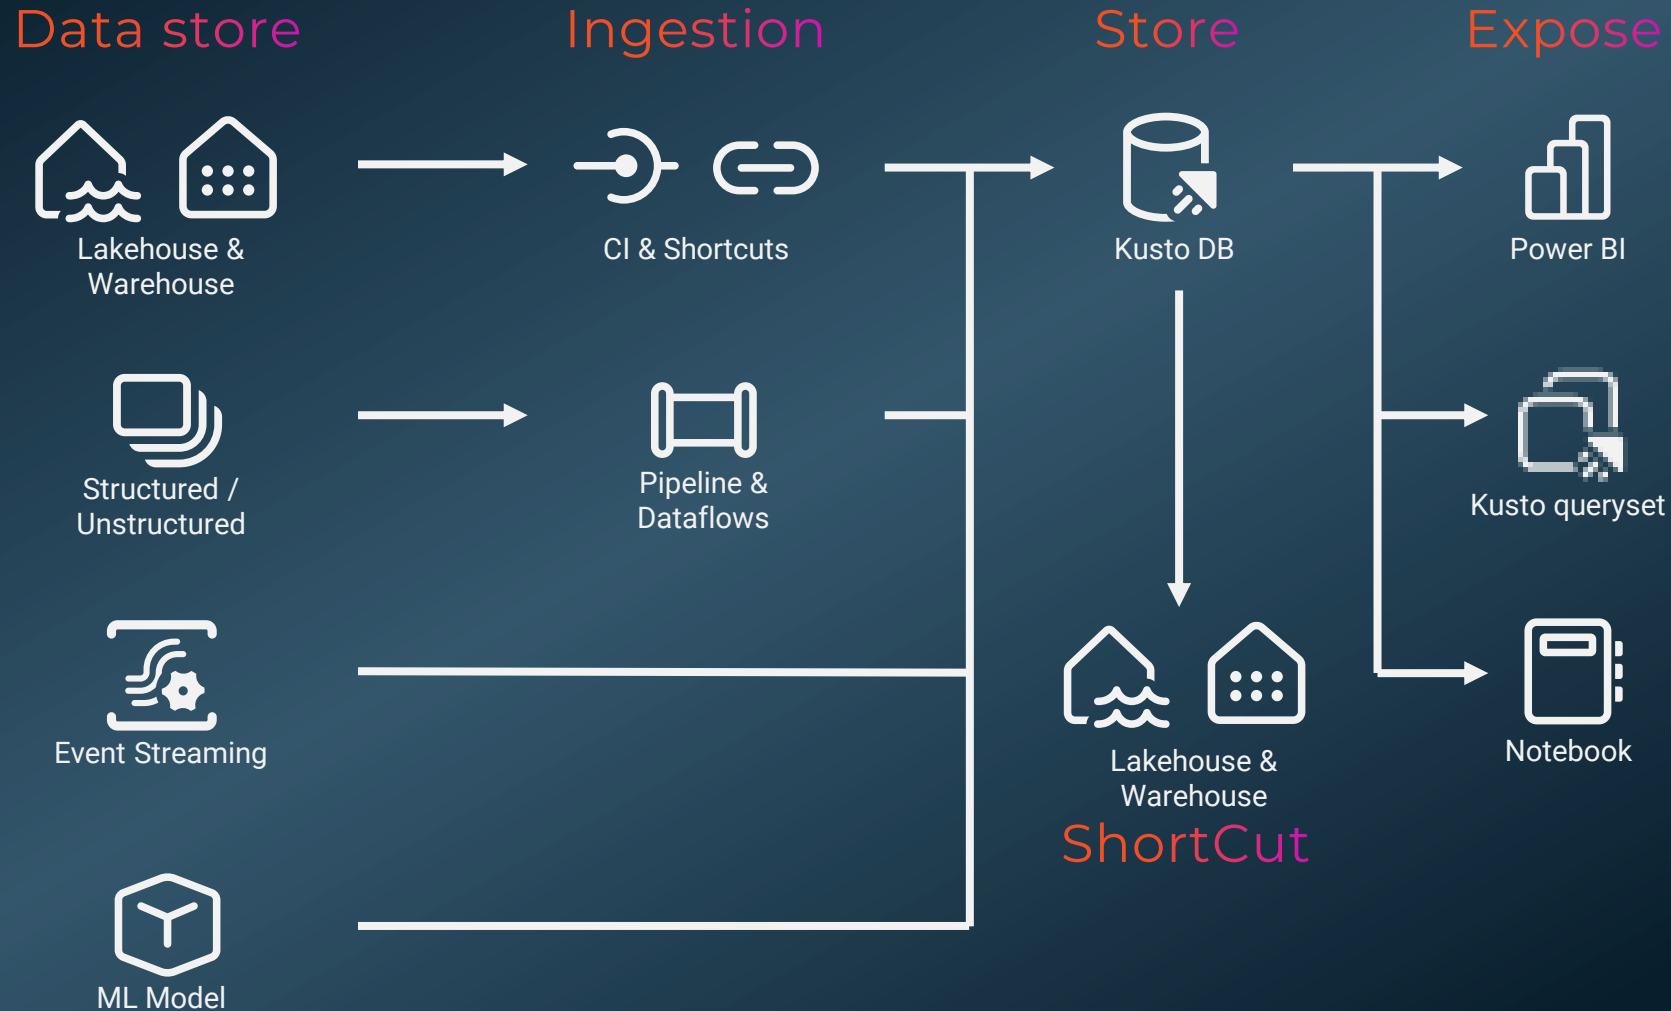
Streaming analytics in
Near-Real-Time

High performance
Low latency
High freshness

Timeseries database

Everything is indexed
and partitioned

Real-Time Analytics



Get started for free

<https://dataexplorer.azure.com/freecluster>

<https://detective.kusto.io>



Kusto in Power BI

Forget everything you know about
query performance vs data types
&
data modelling best practices

Data modelling Kusto in Power BI

- Single table reporting can be a good option, if you can include all columns from dimensions to the table
- M:M relations are hard to avoid, but not a big deal → all queries will be translated to KQL
- All dimensions must be tagged with “IsDimension=true”
- Dimensions can be imported if they are <1 mio rows.
- INTEGER and DECIMAL er slow joins compared to STRING



Harness the Power (BI) of Kusto

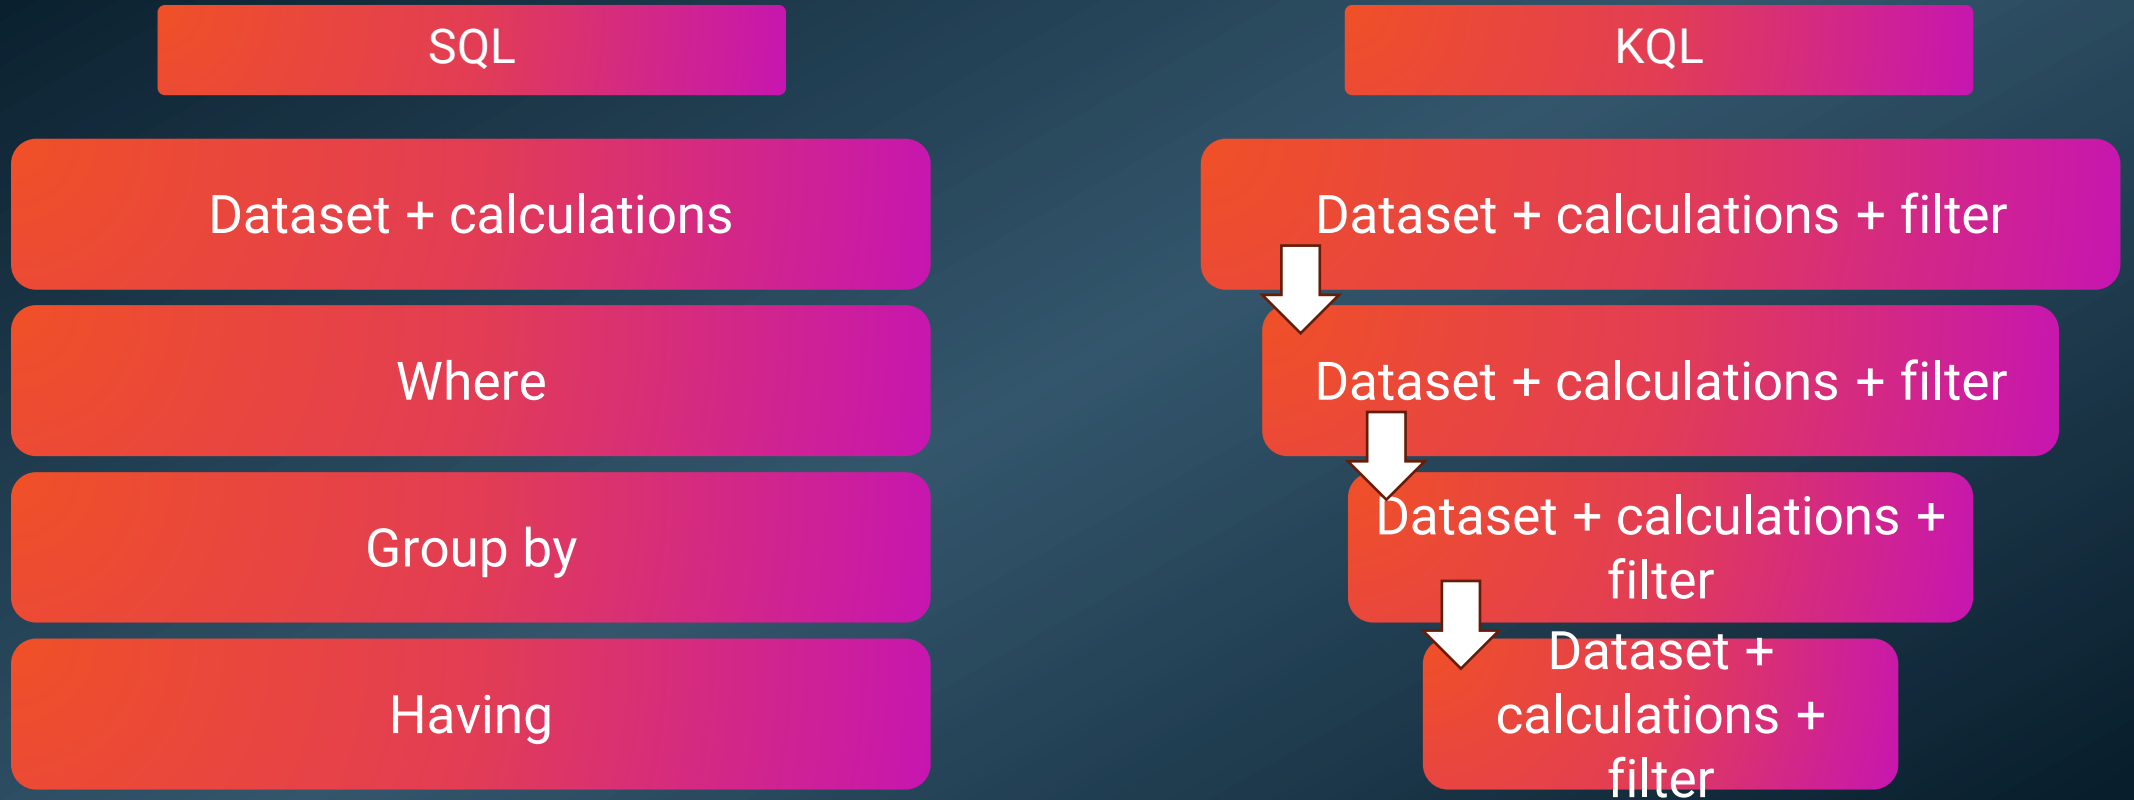
Let Power BI build the KQL

- In Power Query
- Using DAX

Or build a Kusto
function



The language and structure




The language and structure

Using the Notebook feature in Azure Data Studio to demo

Get the database id from your Fabric Kusto cluster

```
print strcat("Database in scope: ", current_database())
```



DEMO

Live coding
(hopefully no demo-ghost 👻)

The language and structure

KQL: Kusto Query Language

SQL

`select * from NYCTaxi`

KQL

`NYCTaxi`

The language and structure

SQL

```
select * from NYCTaxi  
where VendorID = 2
```

KQL

```
NYCTaxi  
| where VendorID == 2
```

The language and structure

SQL

```
select * from NYCTaxi  
where VendorID = 2  
order by passenger_count
```

KQL

```
NYCTaxi  
| where VendorID == 2  
| order by passenger_count
```

The language and structure

SQL

```
select count(*) from  
NYCTaxi
```

KQL

```
NYCTaxi  
| count
```


The language and structure

SQL

```
select
  passenger_count
  ,VendorID
  ,trip_distance
from NYCTaxi
```

KQL

```
NYCTaxi
| project passenger_count, VendorID, trip_distance
```

The language and structure

SQL

```
select
  passenger_count
  ,VendorID
  ,trip_distance
  ,total_amount / passenger_count as AmtPsngr
from NYCTaxi
```

KQL

```
NYCTaxi
| extend AmtPsngr = total_amount / passenger_count
| project passenger_count, VendorID, trip_distance,
  AmtPsngr
```

The language and structure

SQL

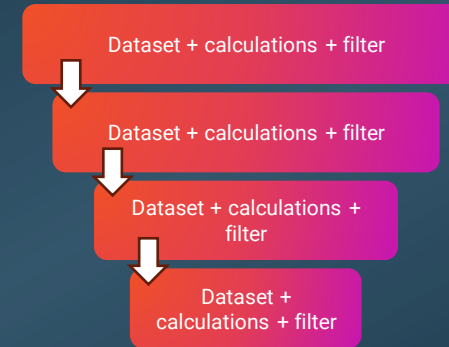
```
select
  sum(passenger_count) as SumPassenger
  ,VendorID
from NYCTaxi
group by VendorID
```

KQL

```
NYCTaxi
| summarize SumPassenger = sum(passenger_count) by
VendorID
```

The language and structure

KQL



NYCTaxi

```
| where passenger_count > 1  
| project passenger_count, total_amount, VendorID, fare_amount  
| extend AmtPsngr = total_amount / passenger_count  
| where AmtPsngr > 10  
| summarize TotalAmount = sum(total_amount), AvgAmtPsngr = avg(AmtPsngr) by VendorID  
| where VendorID <> 1
```


Data discovery and outlier detection

Data discovery is what we've just been through – use select statements and filter your data to find and explore the data given to you.

RENDERING!!

```
NYCTaxi
| where tpep_pickup_datetime between (datetime(2009-01-01)..datetime(2015-01-01))
| extend PickUpdate = startofday(tpep_pickup_datetime)
| summarize SumPsngrCount = sum(passenger_count) by PickUpdate
| project PickUpdate, SumPsngrCount
| render timechart
    with(
        title = "timechart"
        ,xtitle = "Time"
        ,ytitle = "Fares"
    )
```

Data discovery and outlier detection

Outliers `series_outliers()` - [LINK](#)
 `series_decompose()` - [LINK](#)
 `series_decompose_anomalies()` - [LINK](#)
 `series_decompose_forecast()` - [LINK](#)

```
range x from 0 to 364 step 1
| extend t = datetime(2023-01-01) + 1d*x
| extend y = rand() * 10
// generate a sample series with outliers at first day of each month
| extend y = iff(monthofyear(t) != monthofyear(prev(t)), y+20, y)
| summarize t = make_list(t), series = make_list(y)
| extend outliers=series_outliers(series)
| extend pos_anomalies = array_if(series_greater_equals(outliers, 1.5), 1, 0)
| render anomalychart with(xcolumn=t, ycolumns=series, anomalycolumns=pos_anomalies)
```

Functions

Functions in Kusto is equivalent to a stored procedure in the SQL world.

With additional functionality to be able to go outside of the cluster and service and ask for data from a different place in the world.

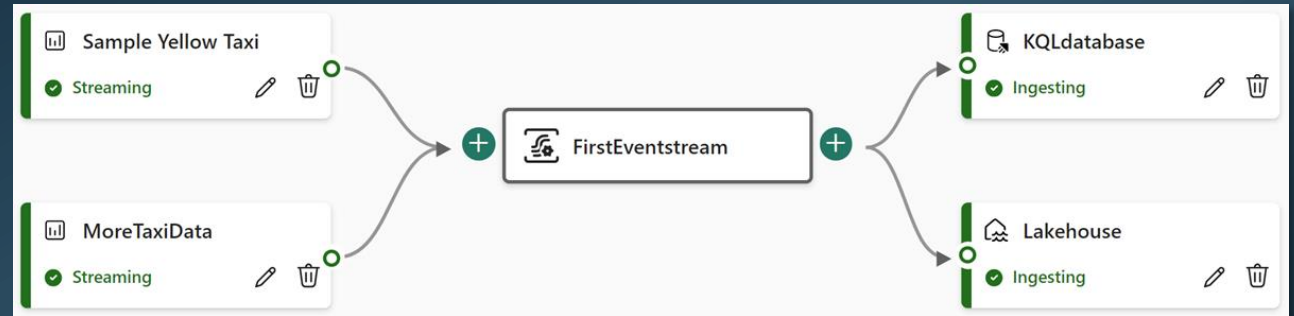
```
.create-or-alter function GetSysLogs(TimeWindow:string , Bucket:string )
{
cluster('help').database('SampleLogs').RawSysLogs
| where timestamp > ago(totimespan(TimeWindow))
| summarize LogCount=count() by name, bin(timestamp, totimespan(Bucket))
| order by timestamp asc
}

// to execute the function
GetSysLogs('5d','1h')
```

Eventstream and Data Activator

Eventstream


The brand-new event stream service, leverages the ability to get data from several sources of streaming data and save it to a wide variety of destinations, including OneLake, KQL databases and Azure services.



Data Activator

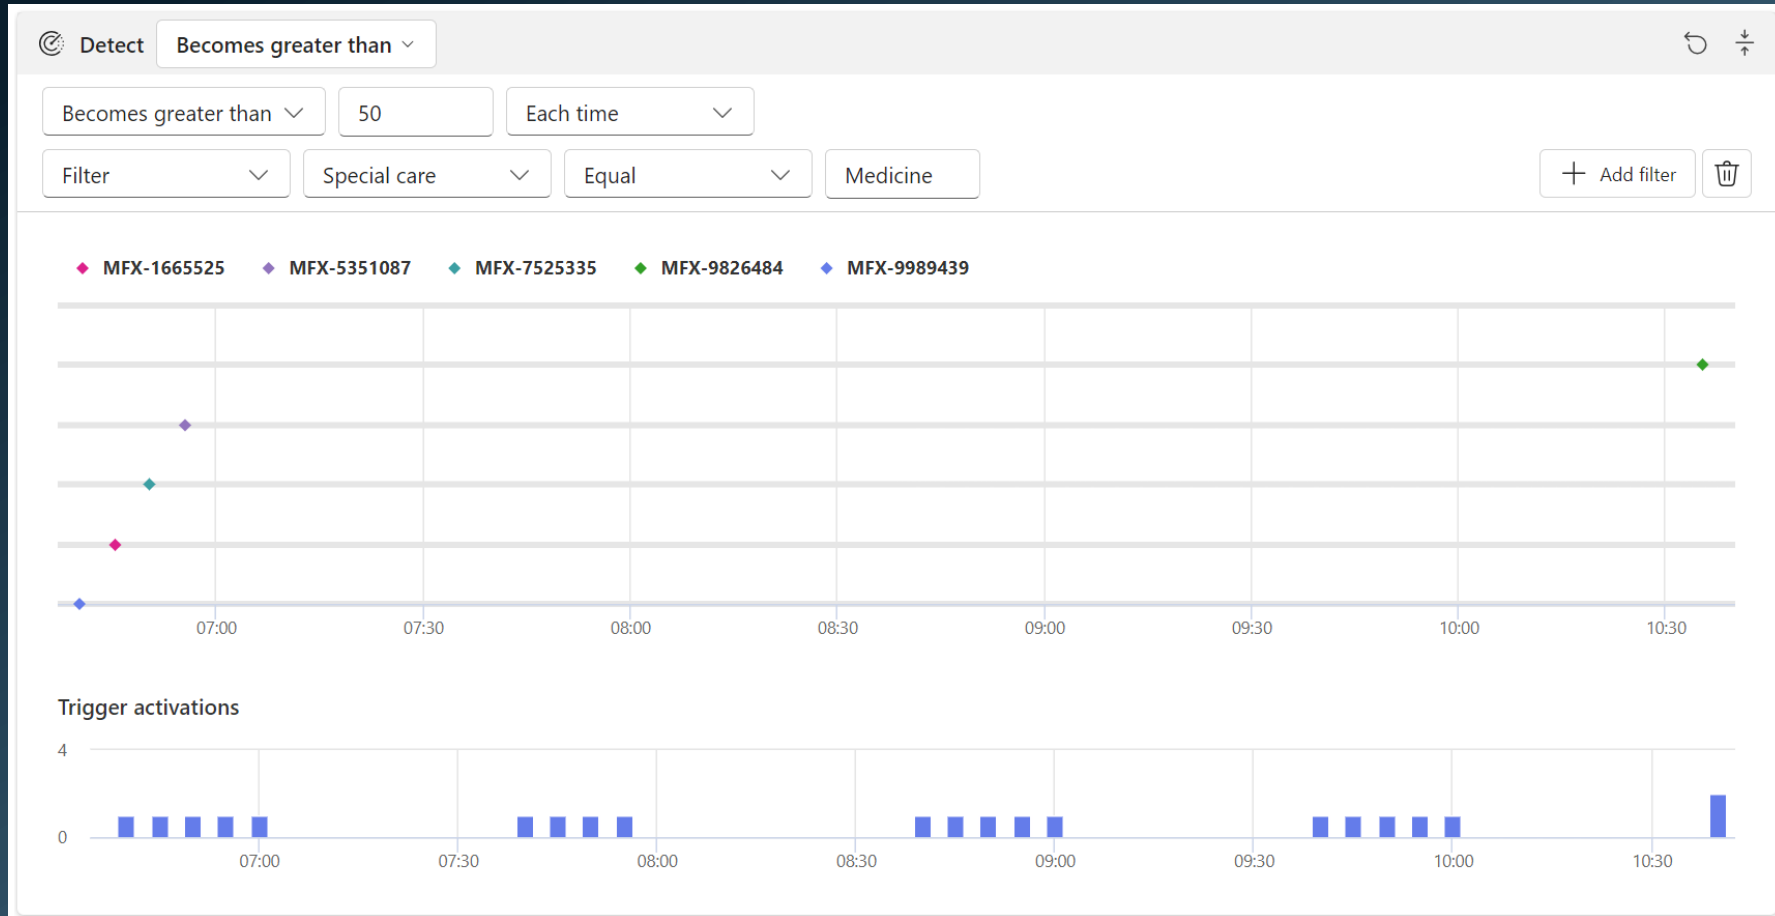
Activeli listens to your data from either the Eventstream service or a Power BI dataset.

Can react to values outside of defined boundaries and, for now, send an e-mail for a Teams message.




Eventstream and Data Activator

Data Activator





Eventstream and Data Activator

Data Activator

 Act

Send email ▾

Send to * ⓘ

BR Brian Børk Rueløkke ✕

Subject * ⓘ

Reflex trigger 'Medicine package is too warm'

Headline * ⓘ

The condition for Reflex trigger 'Medicine package is too warm' has been met

Optional message ⓘ

Additional information ⓘ

Select properties to include in your notification ▾

Limit of 5.

Analysis and reporting



Power BI

Get Data

kusto

All

Azure

Azure Data Explorer (Kusto)

Azure Data Explorer (Kusto)

Advanced options (optional)

Limit query result record number (optional)
Example: 500000

Limit query result data size in Bytes (optional)
Example: 67108864

Disable result-set truncation (optional)
Example: false

Additional Set Statements (separated by semicolons... (optional)
Example: set query_datascope=hotcache;

Data Connectivity mode ⓘ

☐ Import

☒ DirectQuery

OK Cancel

DEMO

Live coding
(hopefully no demo-ghost 👻)

Harness the Power (BI) of Kusto

```
.create-or-alter function GetSysLogs(TimeWindow:string , Bucket:string )
{
cluster('help').database('SampleLogs').RawSysLogs
| where timestamp > ago(totimespan(TimeWindow))
| summarize LogCount=count() by name, bin(timestamp,
totimespan(Bucket))
| order by timestamp asc
}
```

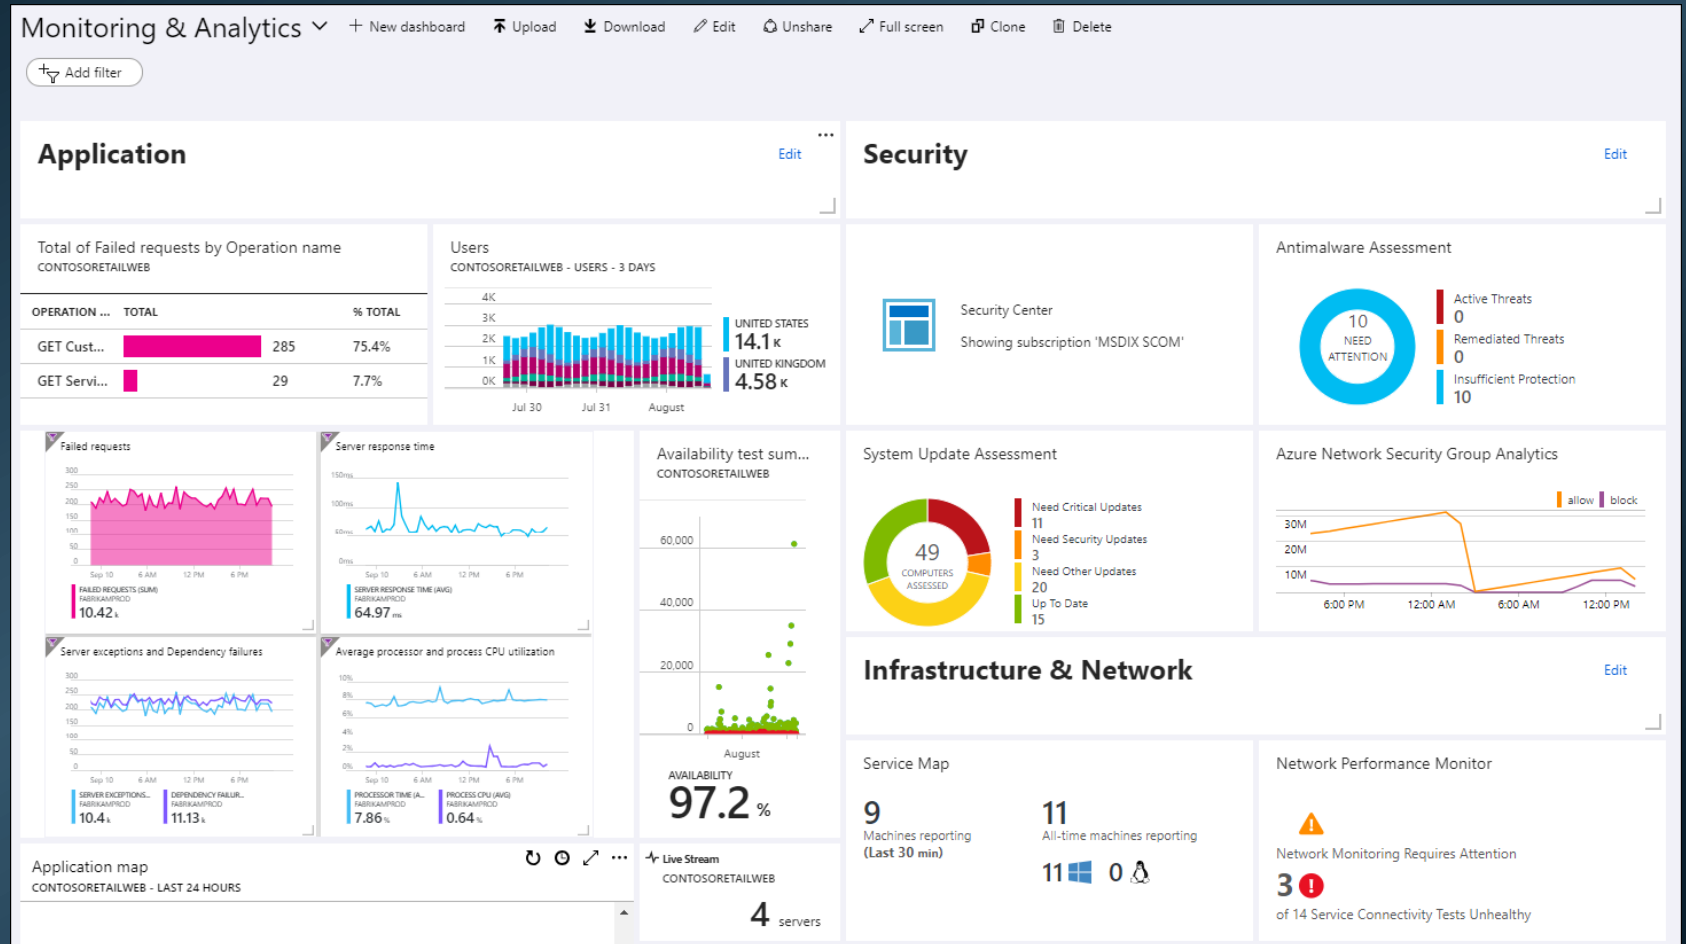
```
GetSysLogs('5d','1h')
```


Analysis and reporting



Real-Time
Analytics

Dashboards in RTA - planned - to come...



Thank you

Connect with me at:

 <https://linkedin.com/in/brianbonk>
 <https://brianbonk.dk>
 <https://github.com/brianbonk>



